



The Containers and Cloud-Native Roadshow Developer Track Lab Guide

A hands-on experience for Ops and Dev professionals



- \$ MAN WHOAMI?
- WHAT I WANT TO LEARN
- FUN FACT (OPTIONAL)

Laine Vyvyan



Josh Smith



DEVELOPER TRACK MODULES

1 OPTIMIZING EXISTING APPLICATIONS

Migrate an existing monolithic Java application from a legacy platform to Red Hat.

Modernize by incrementally refactoring to microservices architecture and modern Java platform

3 CONTROL CLOUD NATIVE APPS WITH SERVICE MESH

Gain a deep understanding of app behavior through service mesh instrumentation and visualization

2 DEBUGGING, MONITORING AND CONTINUOUS DELIVERY

Debug, instrument and monitor a modern microservice application.

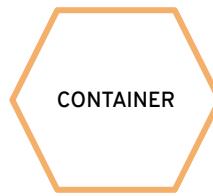
Deploy continuously using Pipelines

4 ADVANCED CLOUD NATIVE WITH EVENT-DRIVEN SERVERLESS

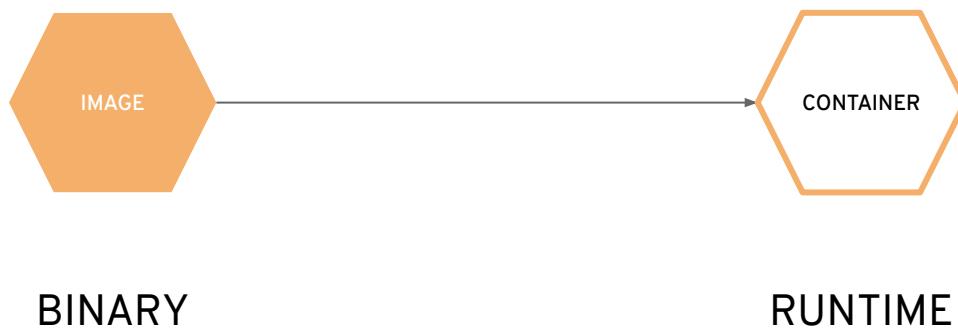
Dynamically respond to events and scale applications using powerful Kubernetes constructs

OpenShift Concepts

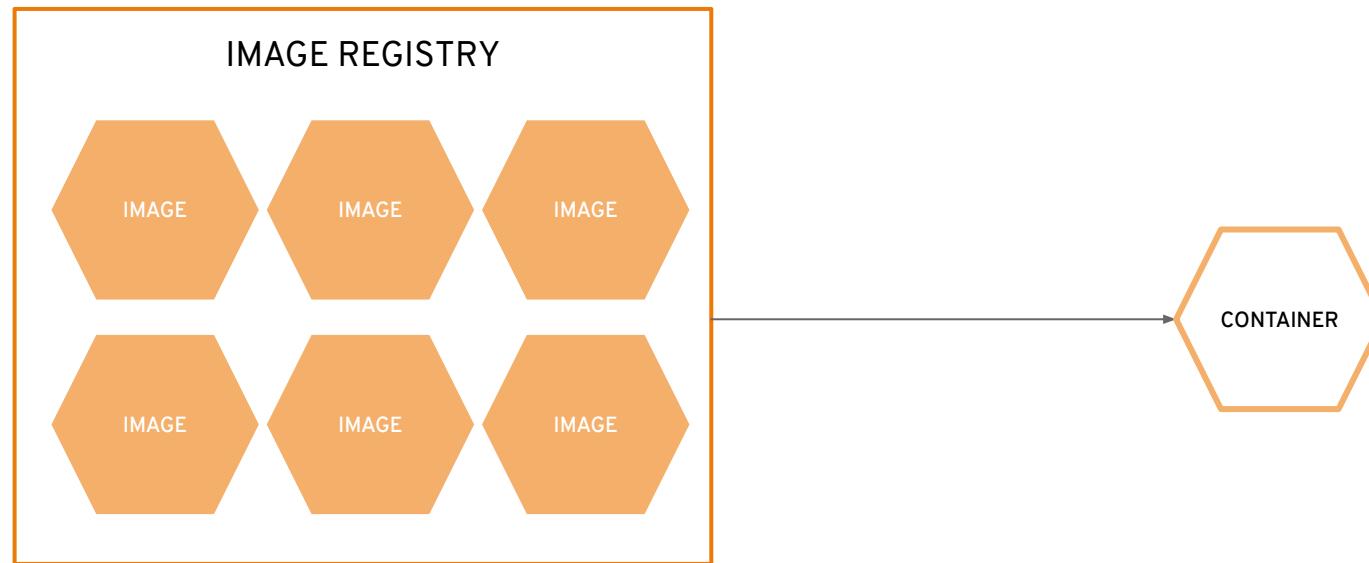
a container is the smallest compute unit



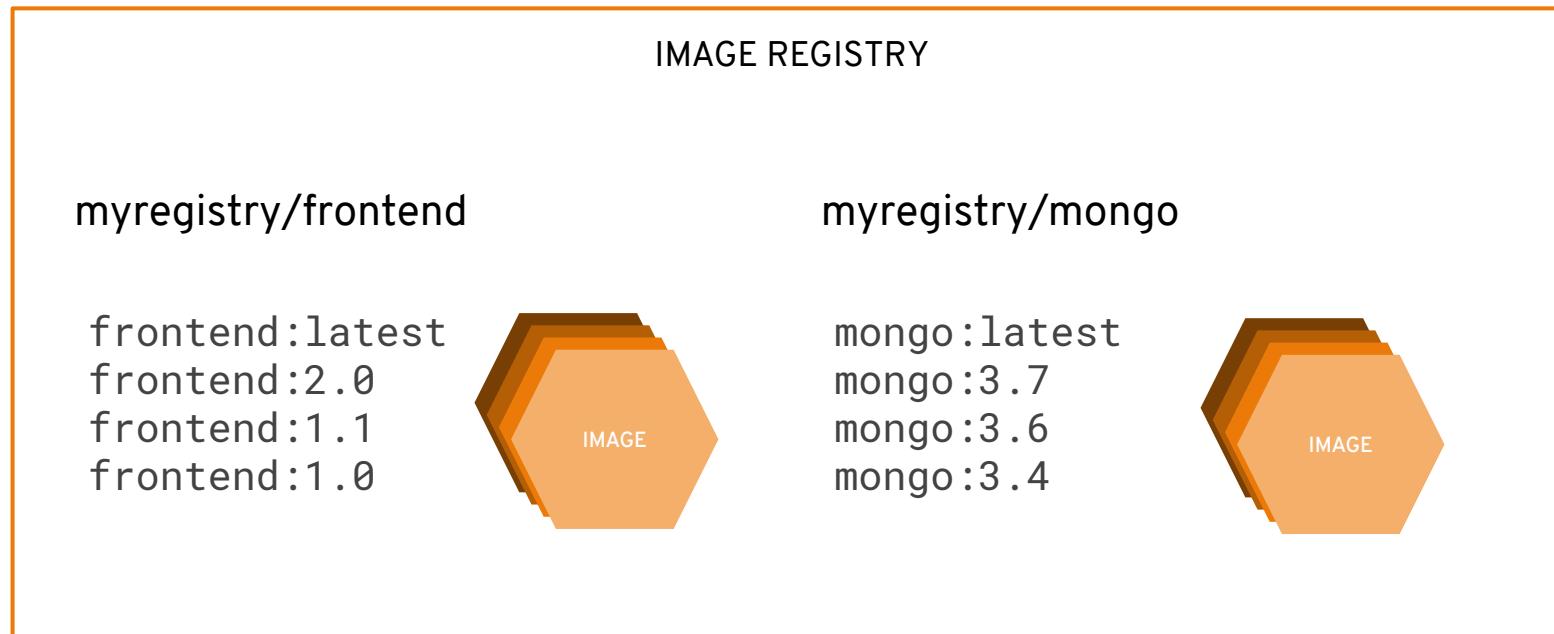
containers are created from container images



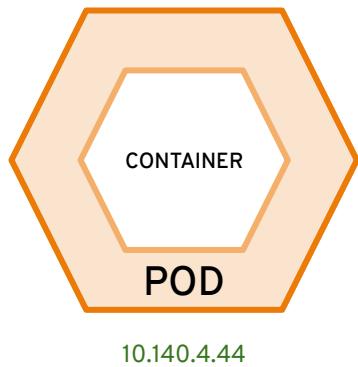
container images are stored in an image registry



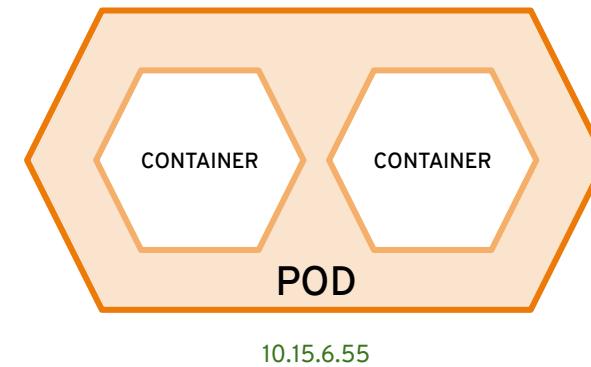
an image repository contains all versions of an image in the image registry



containers are wrapped in pods which are units of deployment and management

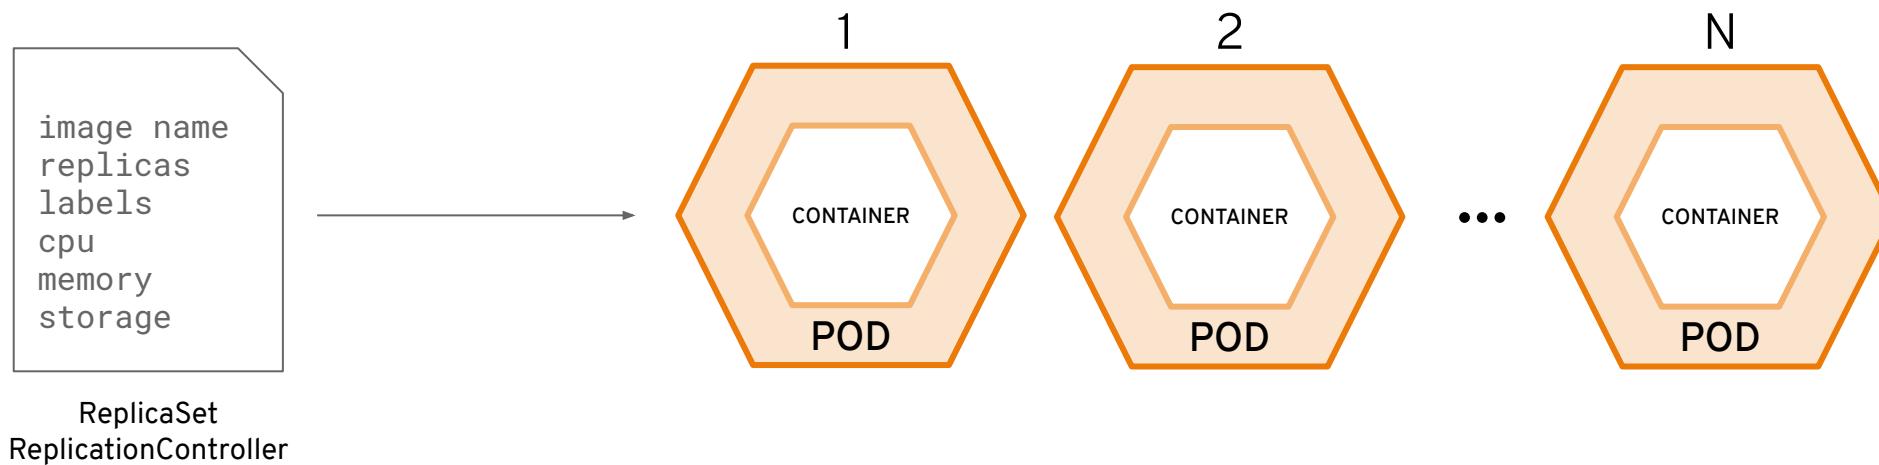


10.140.4.44

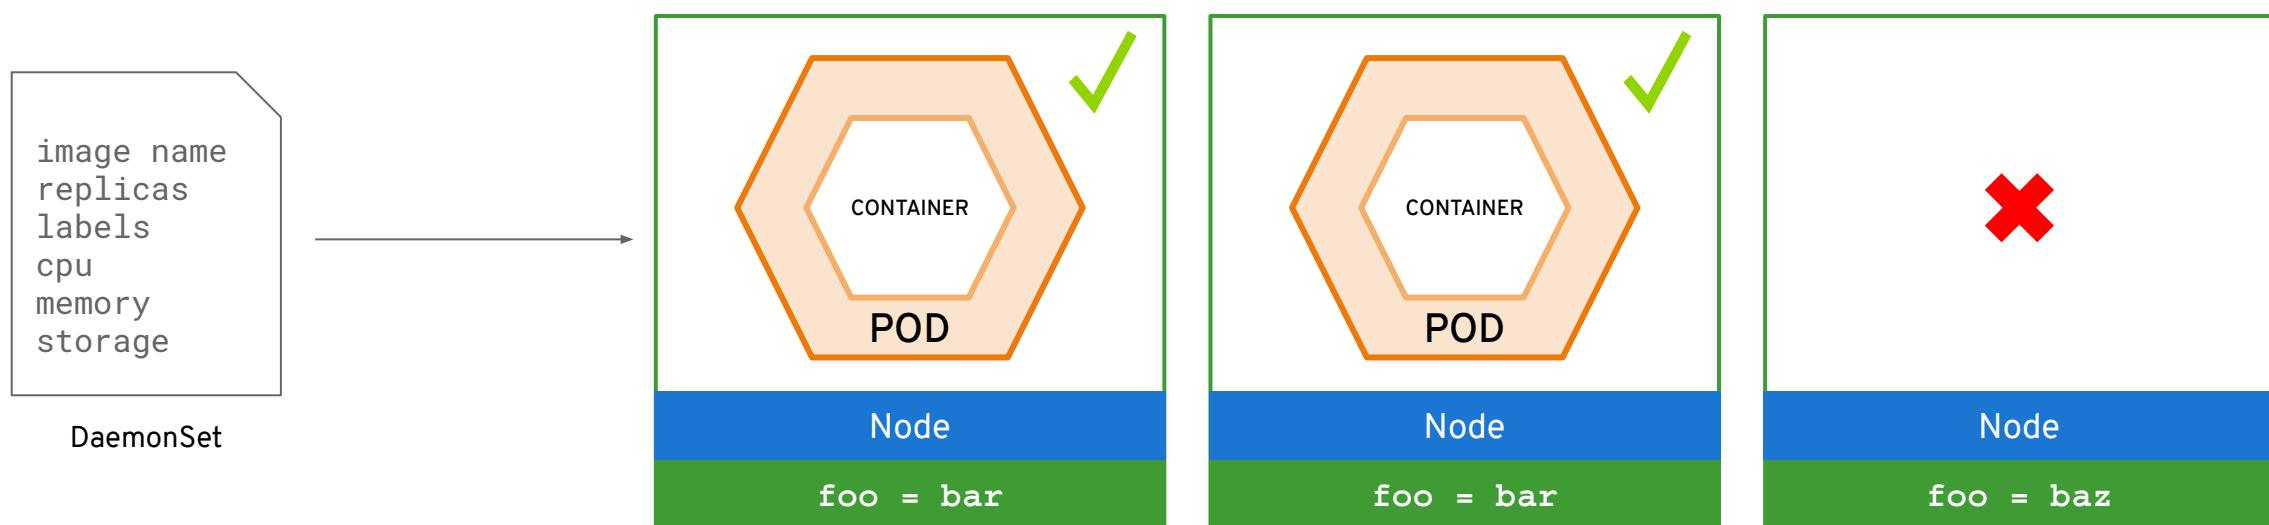


10.15.6.55

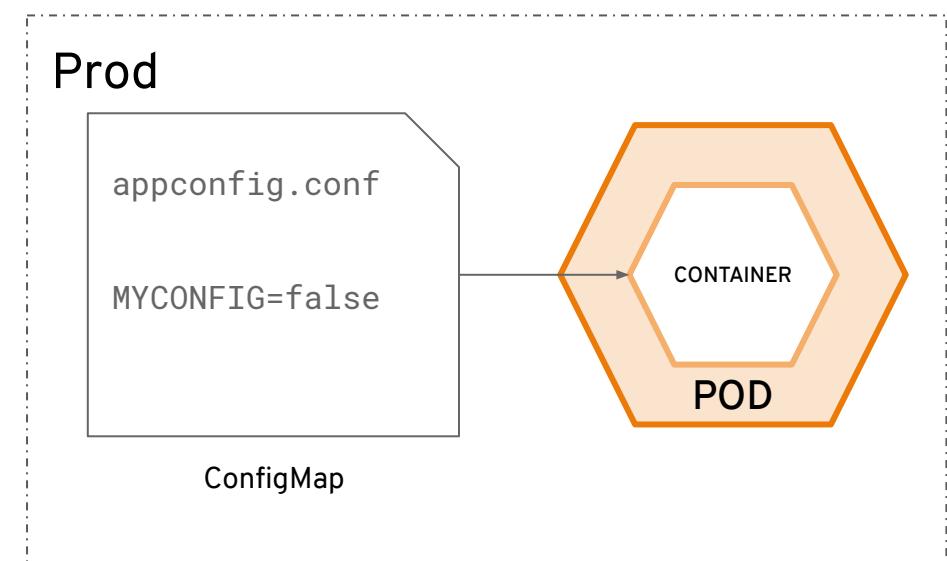
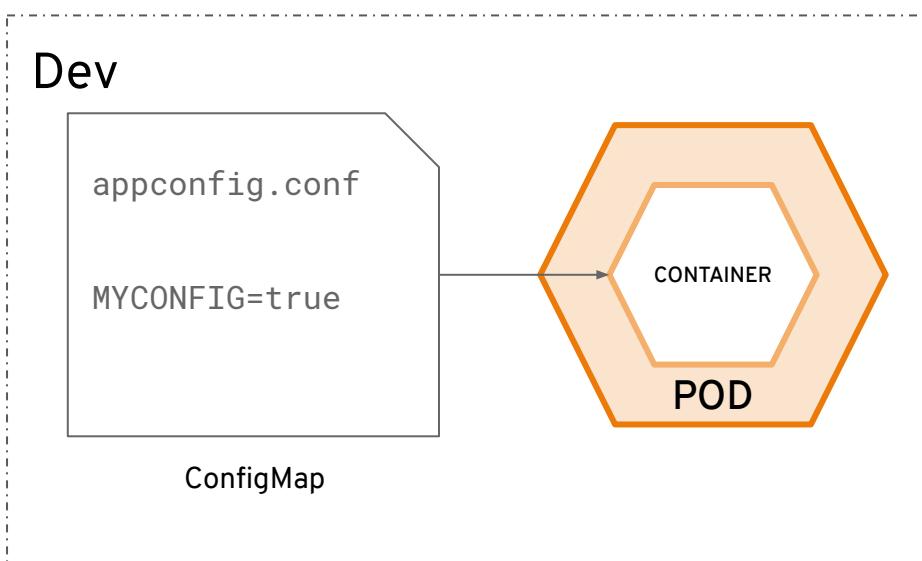
ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



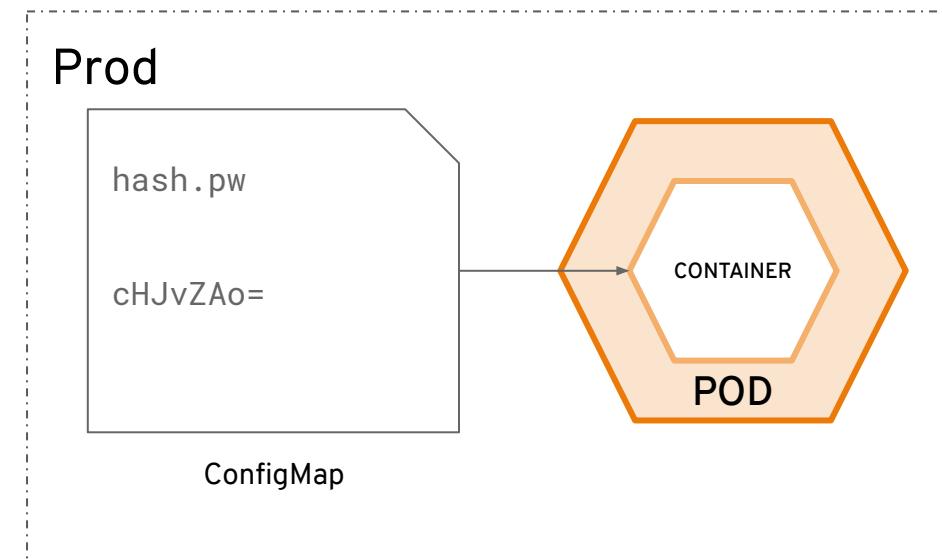
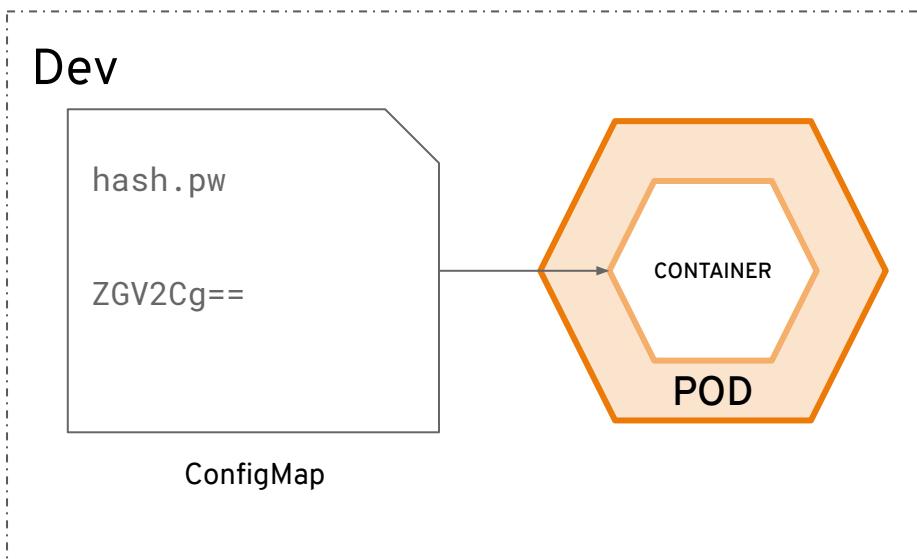
a daemonset ensures that all (or some) nodes run a copy of a pod



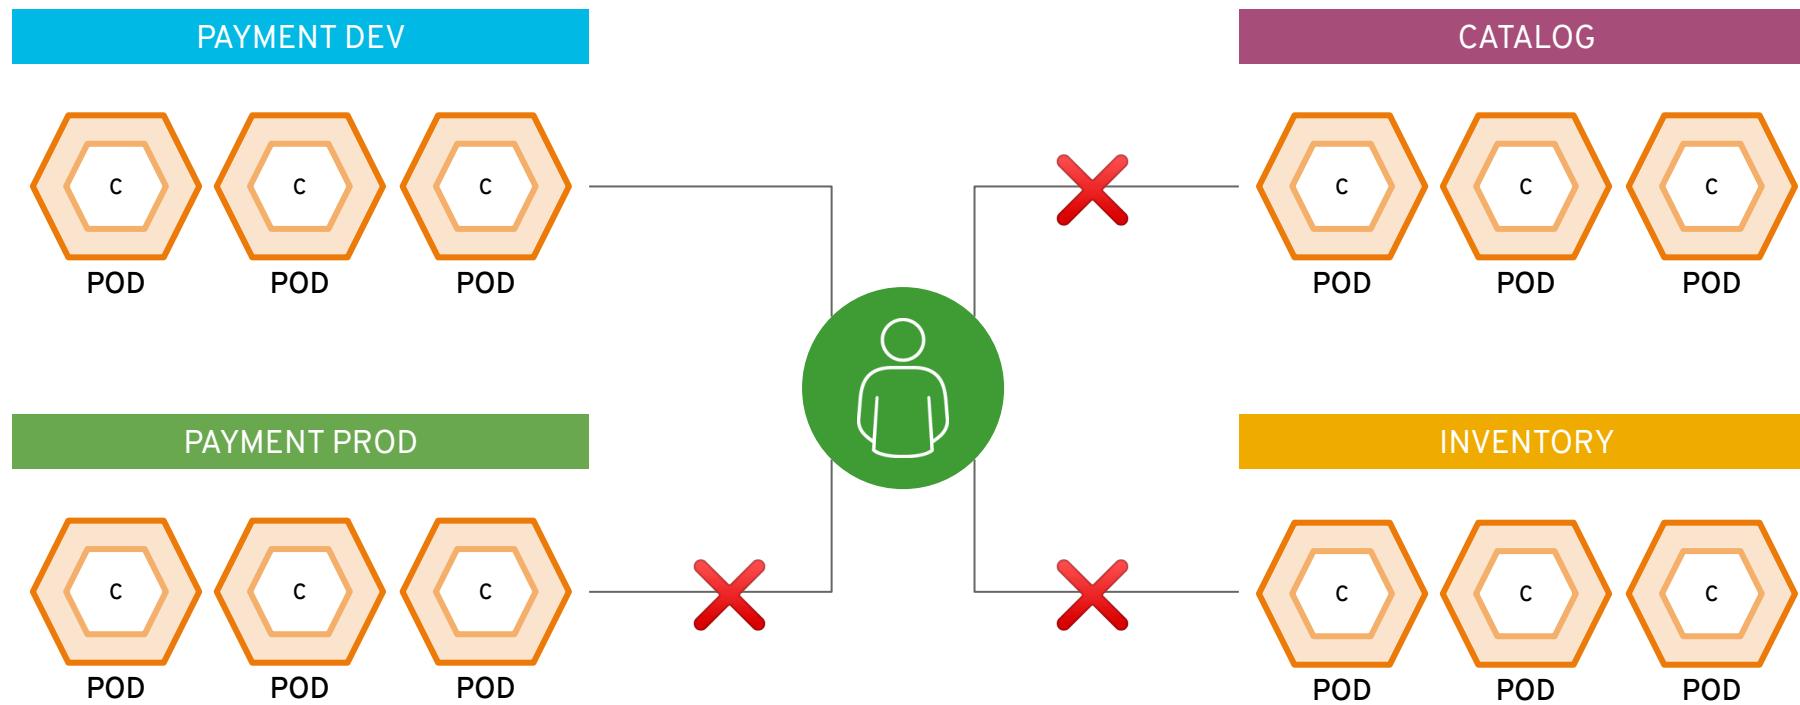
configmaps allow you to decouple configuration artifacts from image content



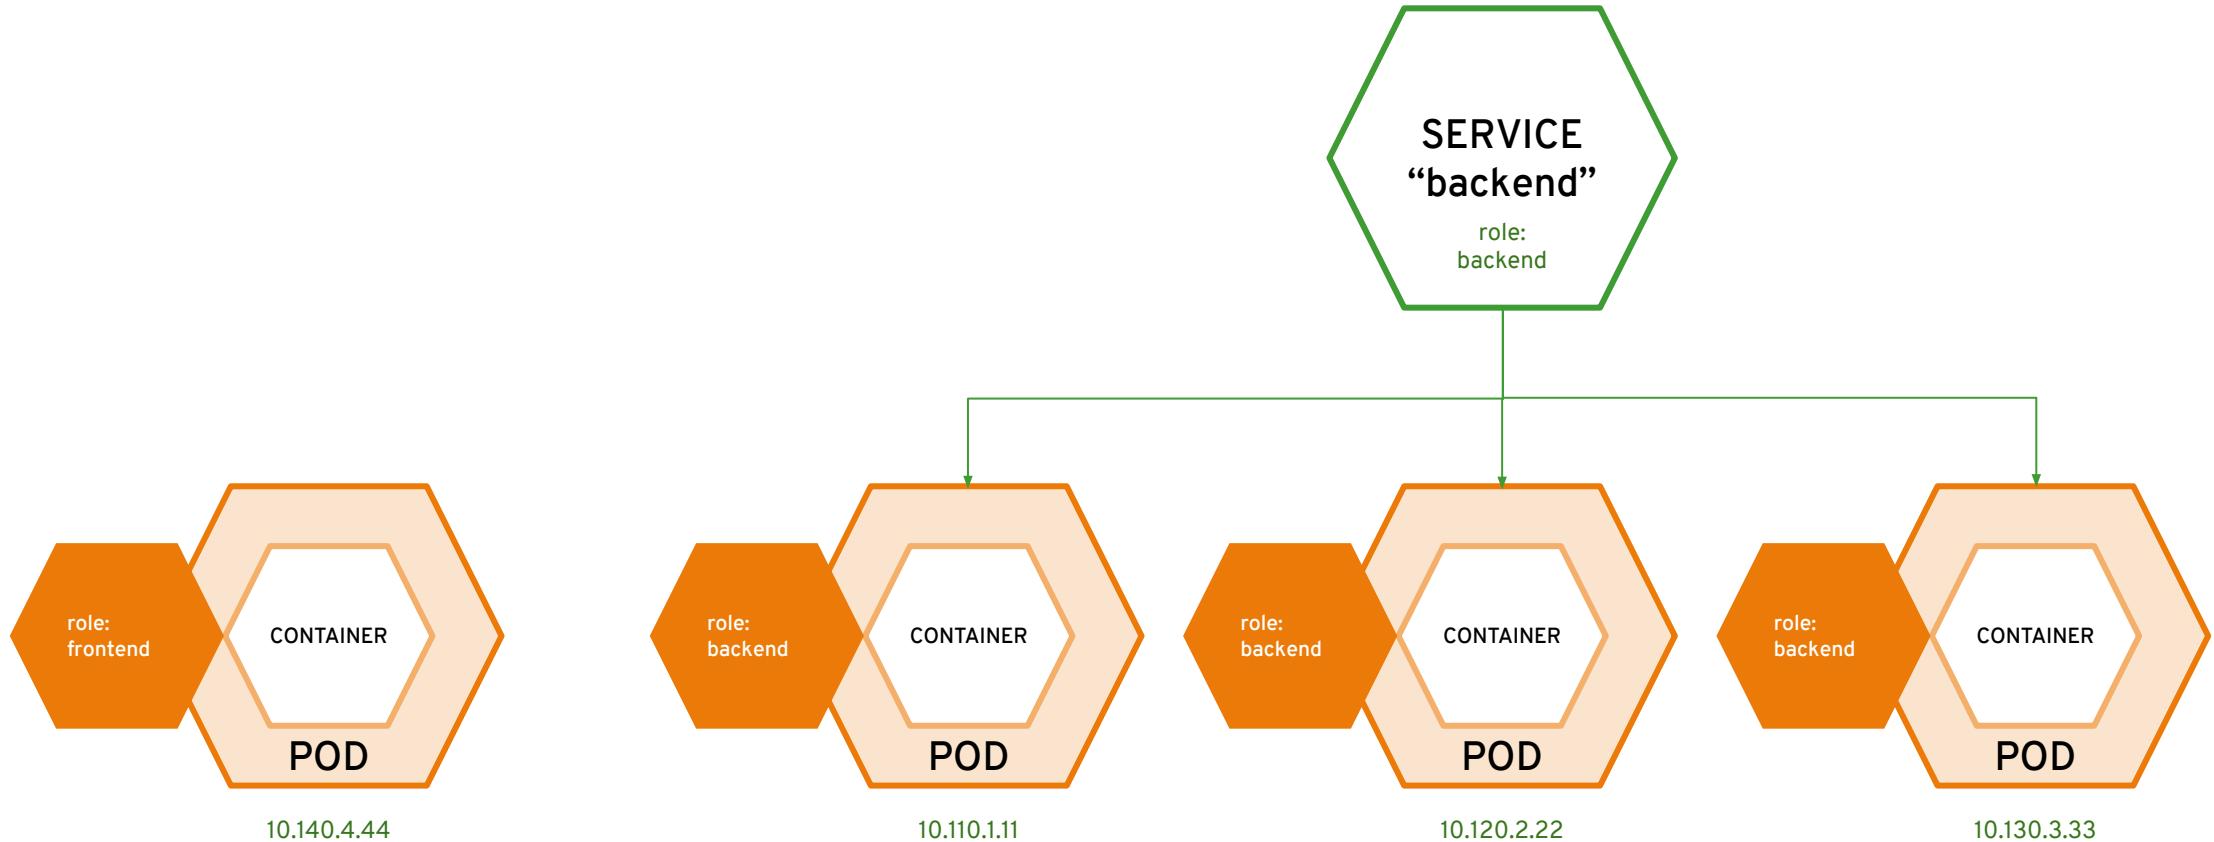
secrets provide a mechanism to hold sensitive information such as passwords



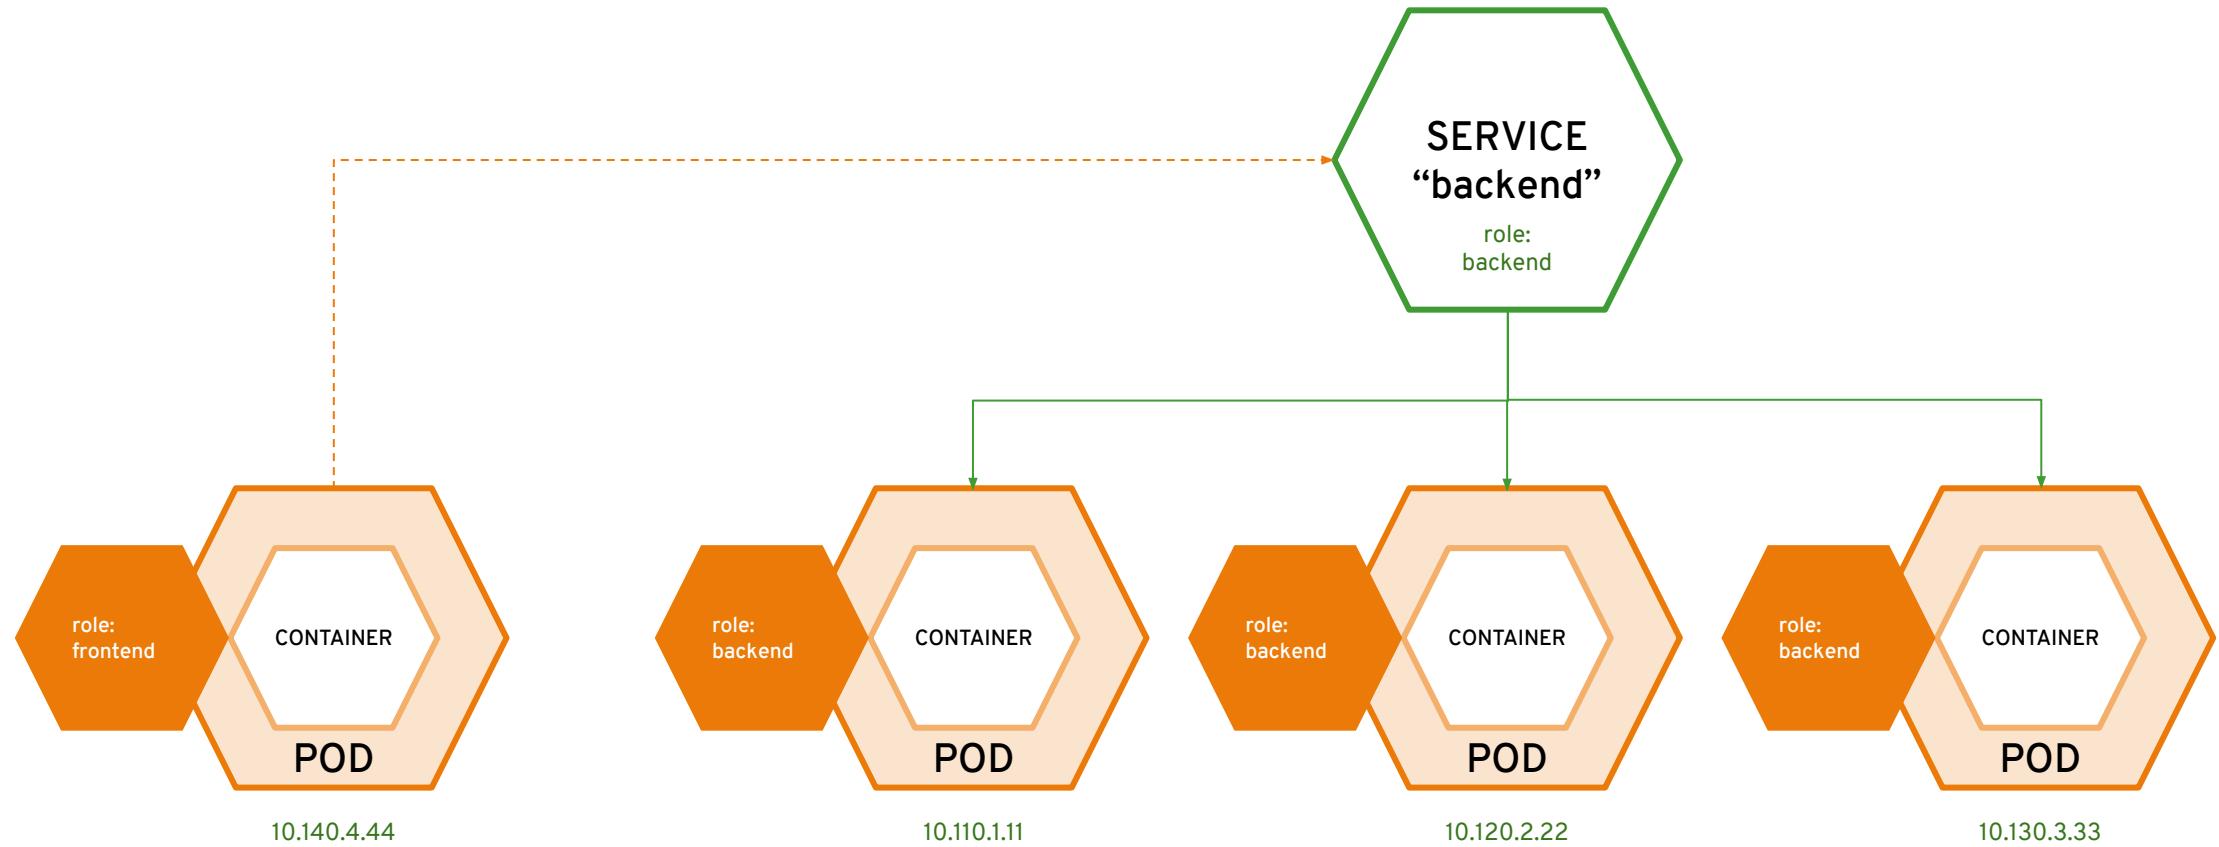
projects isolate apps across environments,
teams, groups and departments



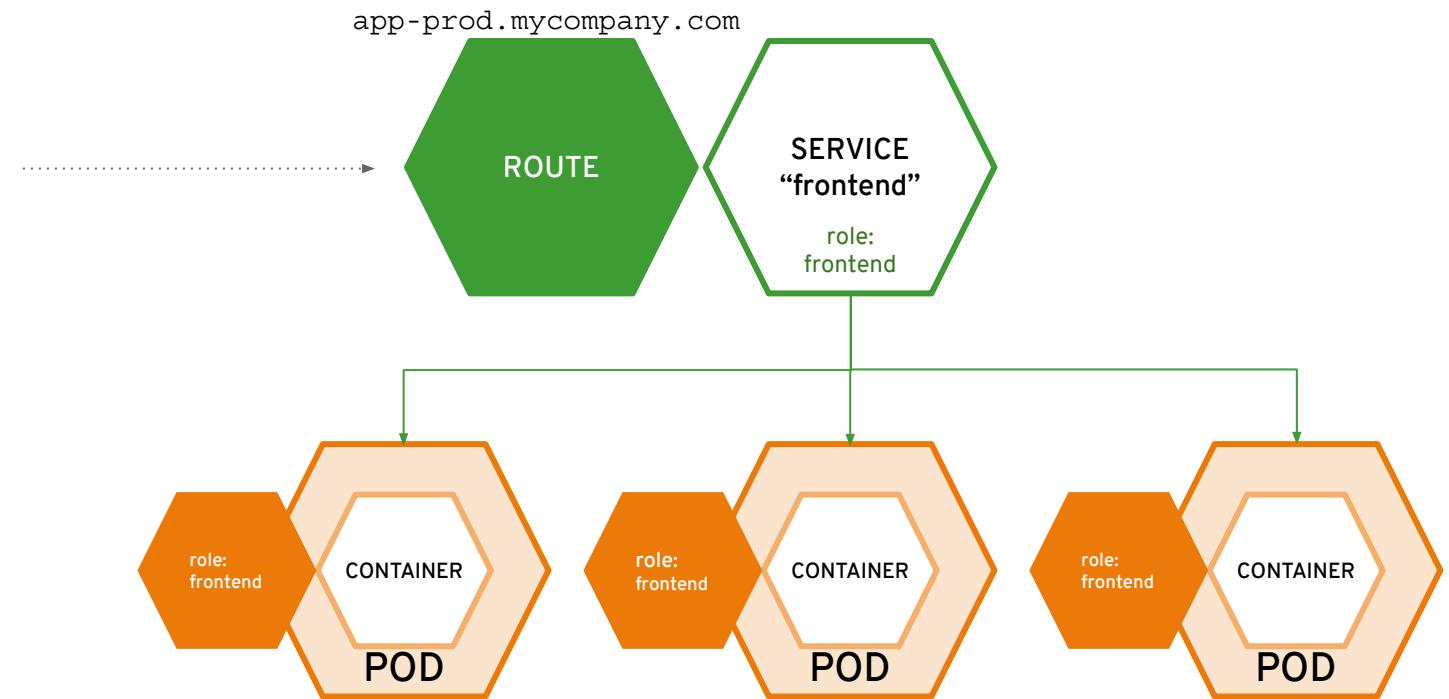
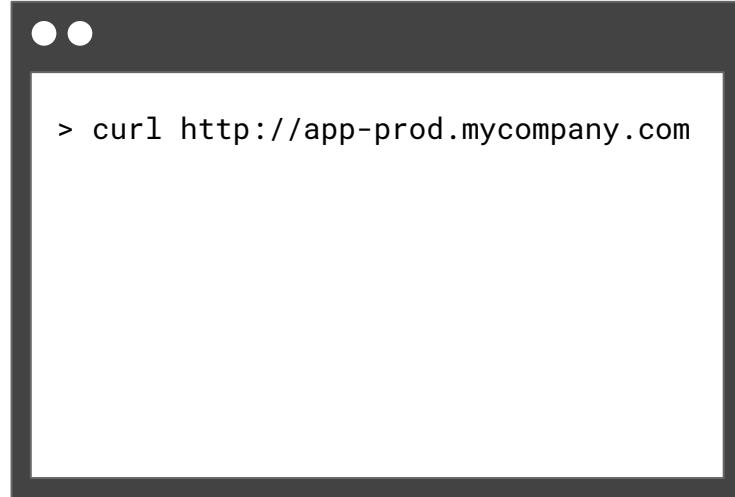
services provide internal load-balancing and service discovery across pods

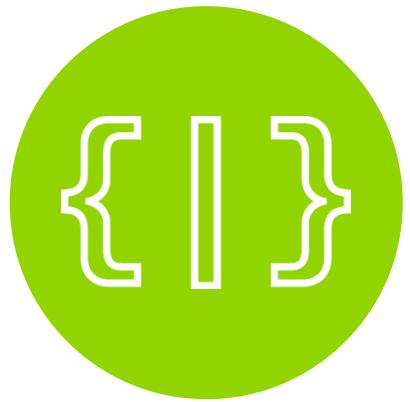


apps can talk to each other via services

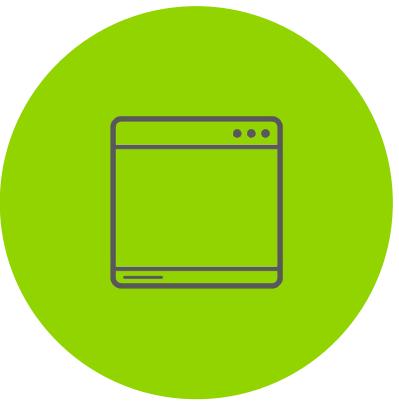


routes make services accessible to clients outside the environment via real-world urls

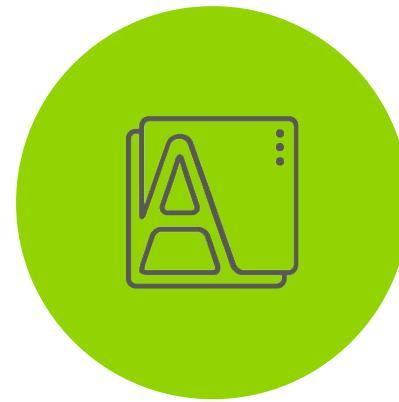




**DEPLOY YOUR
SOURCE CODE**



**DEPLOY YOUR
APP BINARY**



**DEPLOY YOUR
CONTAINER IMAGE**

BUILD AND DEPLOY | Source-to-Image (S2I) for building and deploying from code

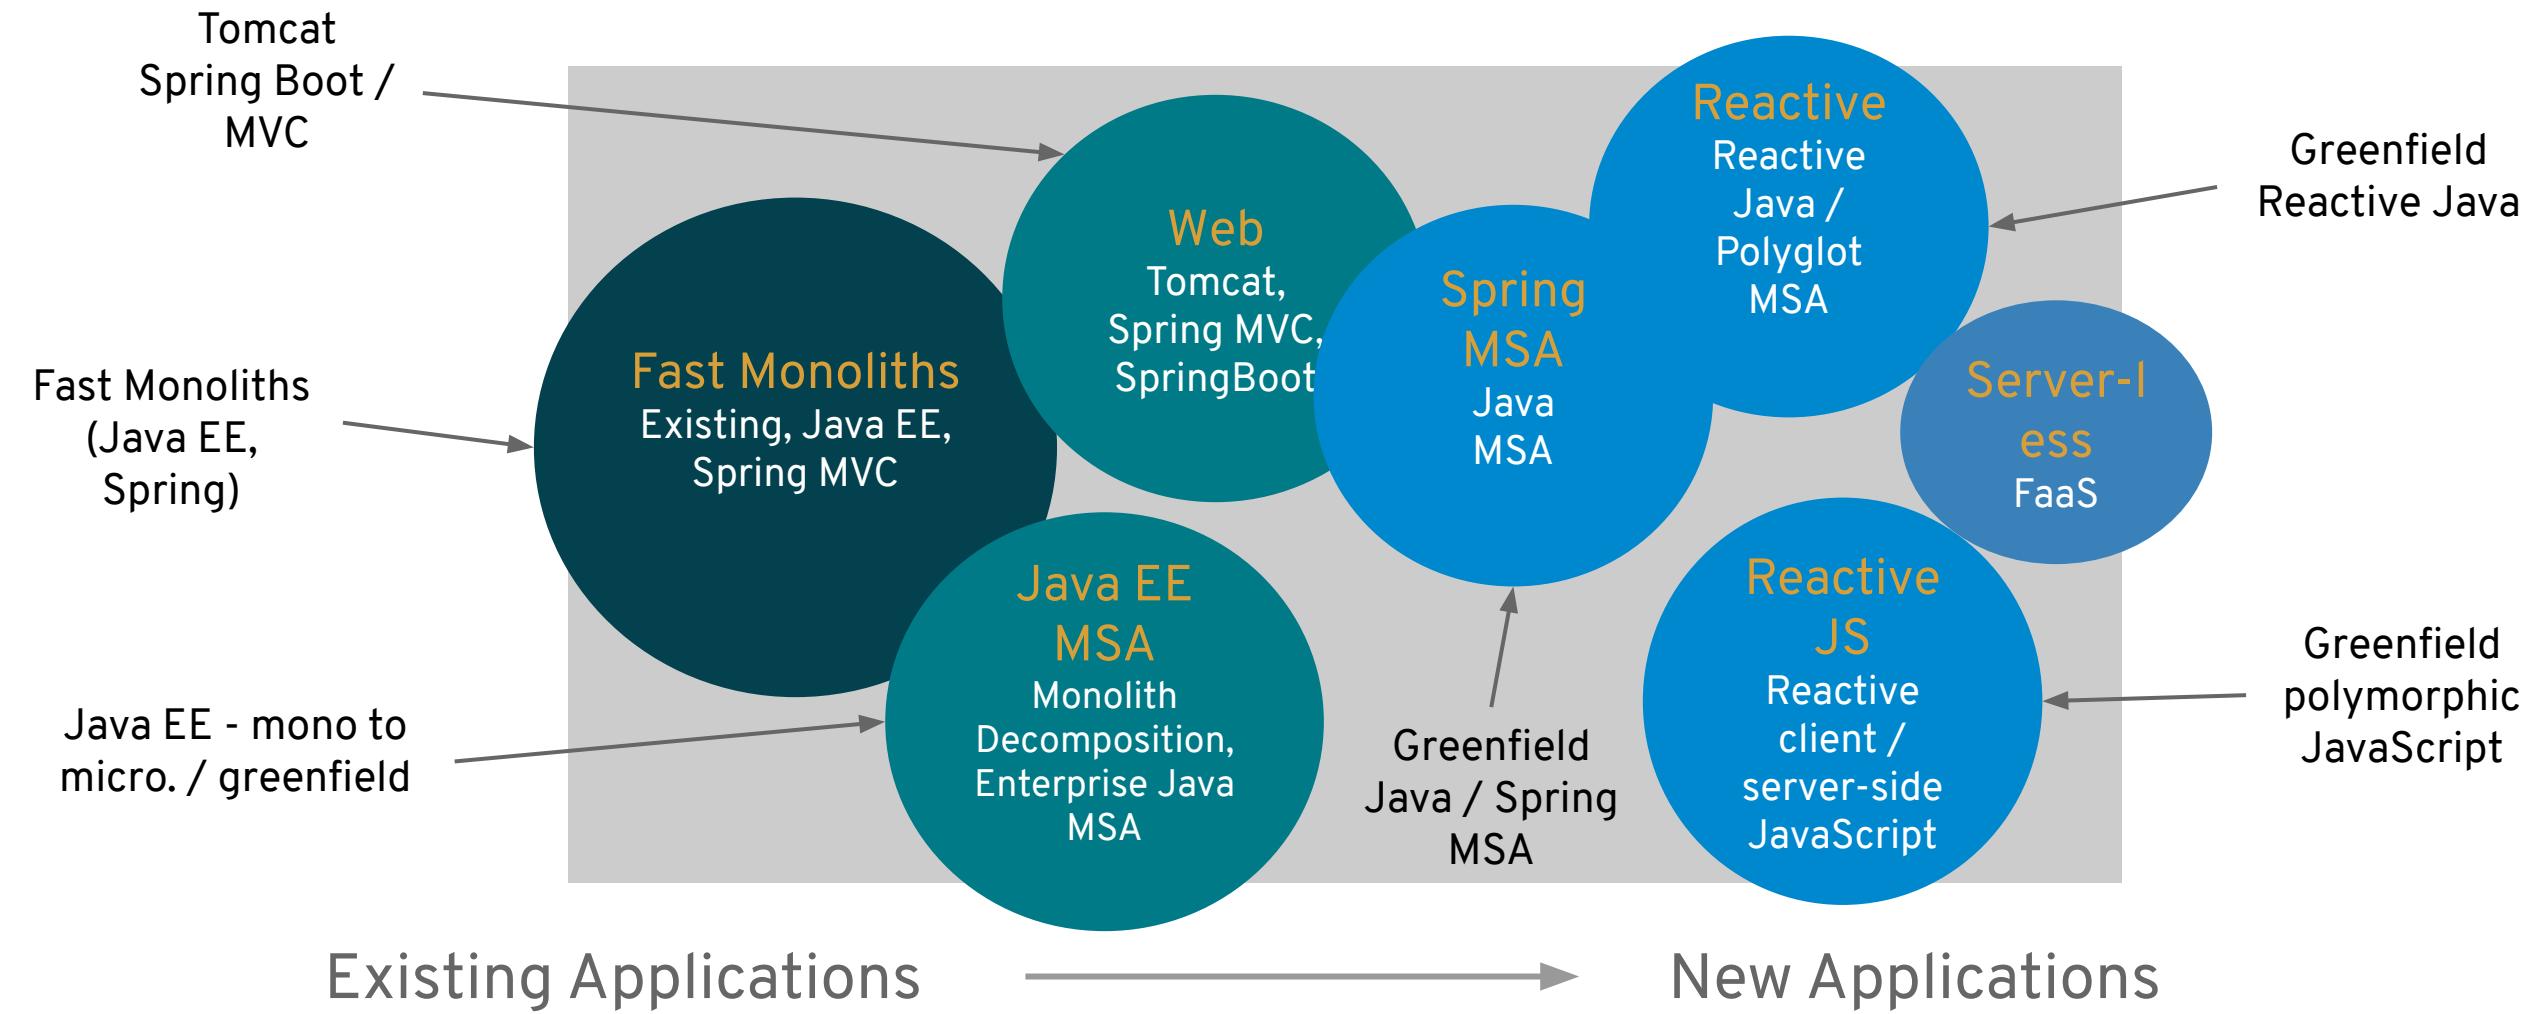


Module 1: Optimizing Existing Applications

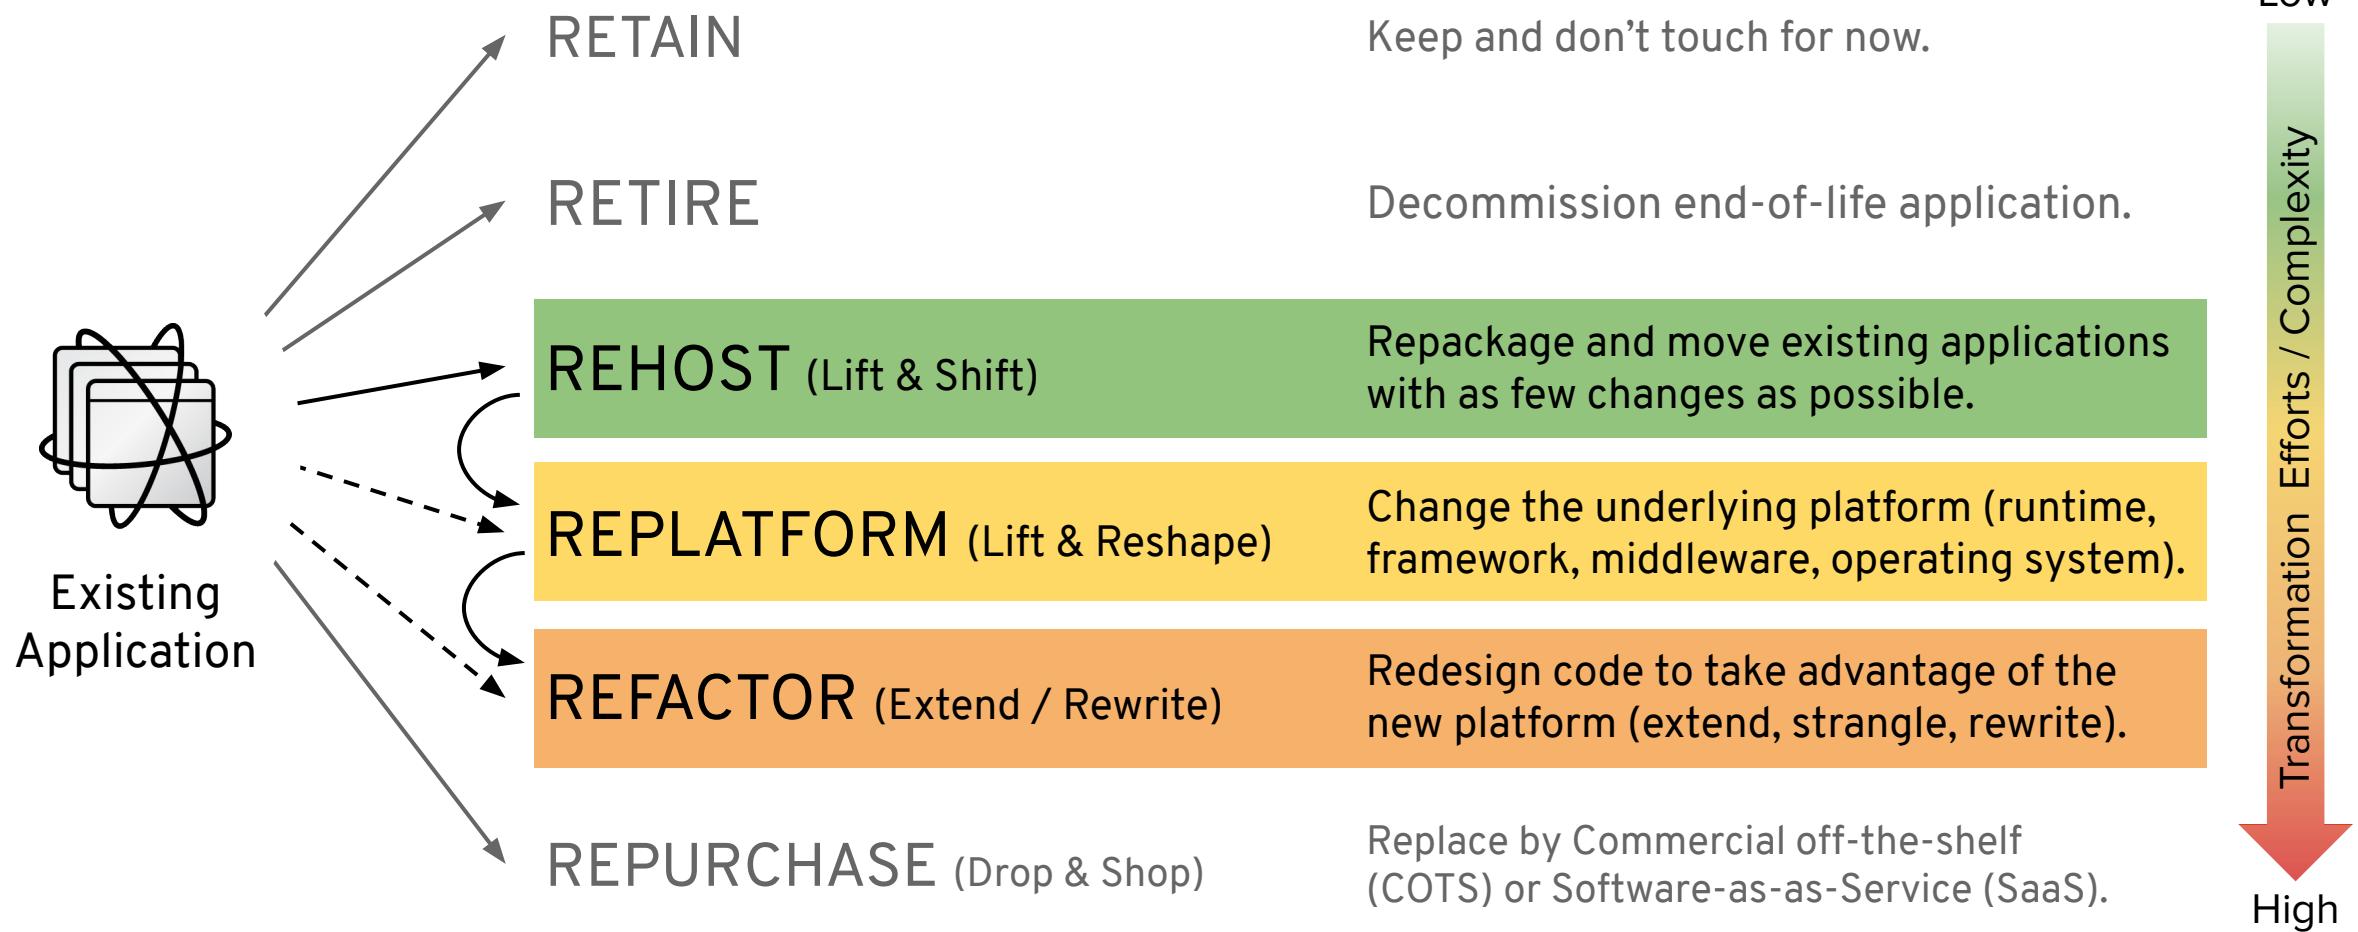
- 1. Optimizing Existing Applications
- 2. Debugging, Monitoring, and Continuous Delivery
- 3. Control Cloud Native Apps with Service Mesh
- 4. Advanced Cloud Native with Event-Driven Serverless



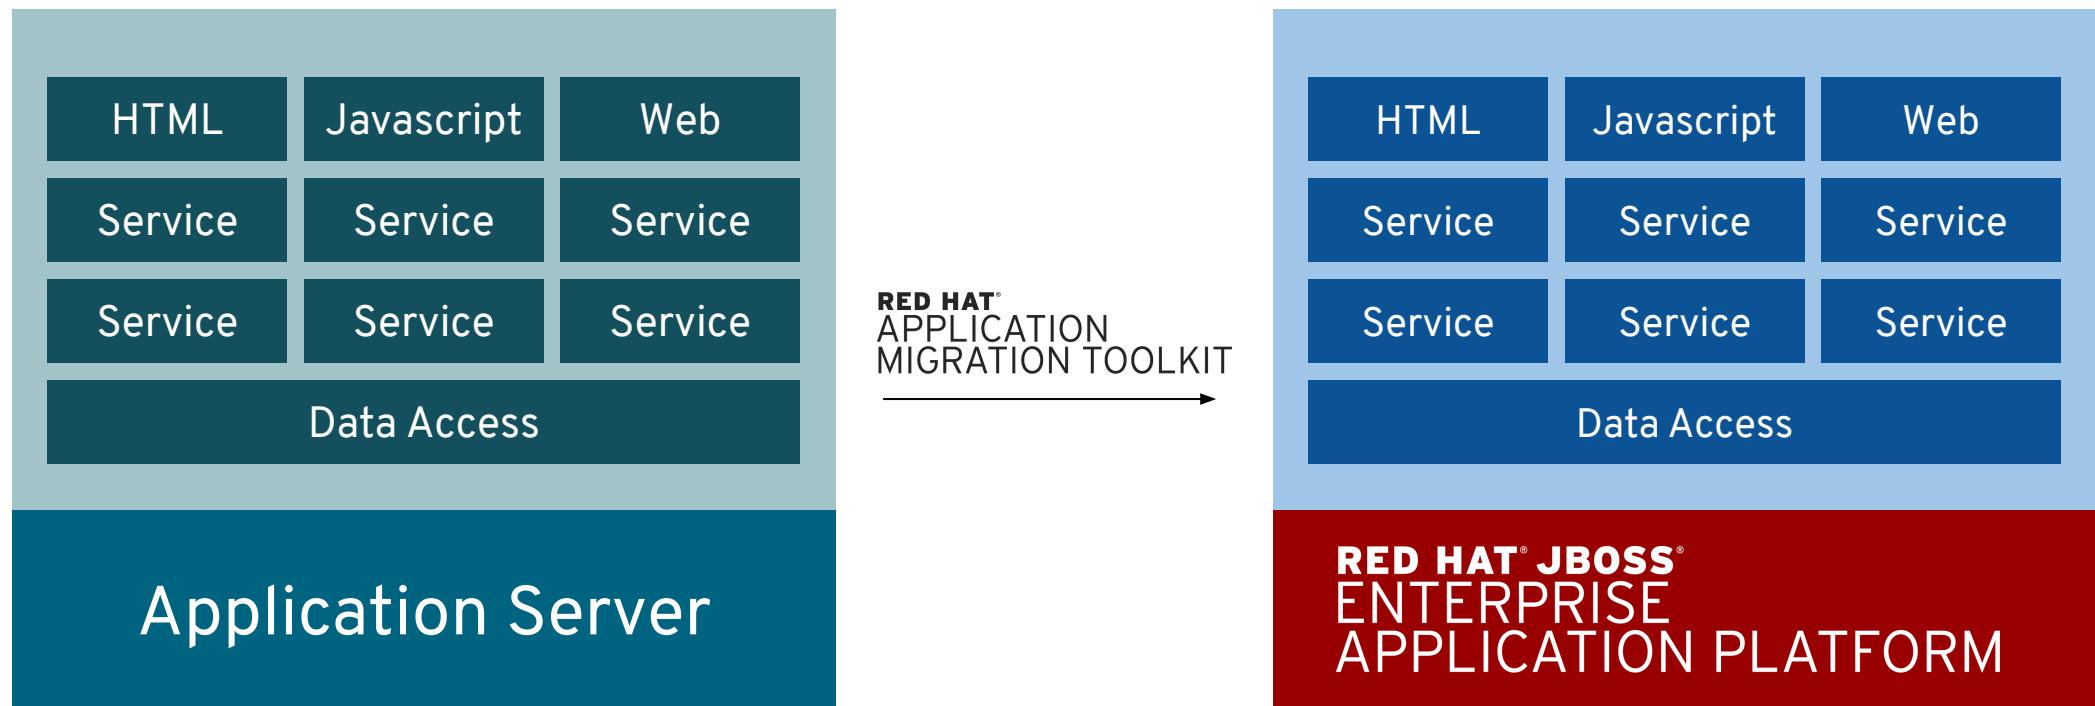
THE SPECTRUM OF ENTERPRISE APPS



OPTIONS FOR CONTAINERIZATION



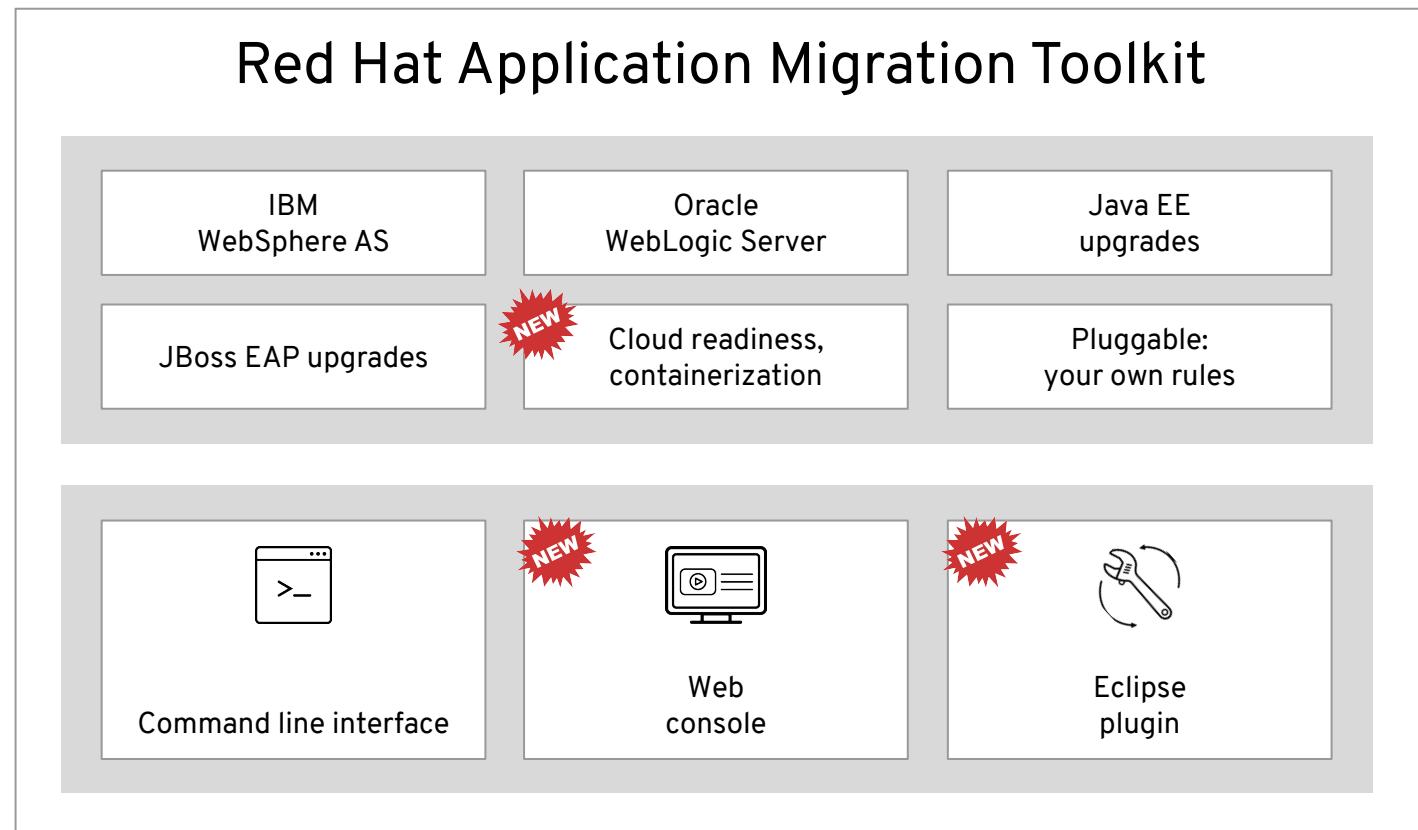
LIFT-AND-SHIFT MONOLITH TO CLOUD



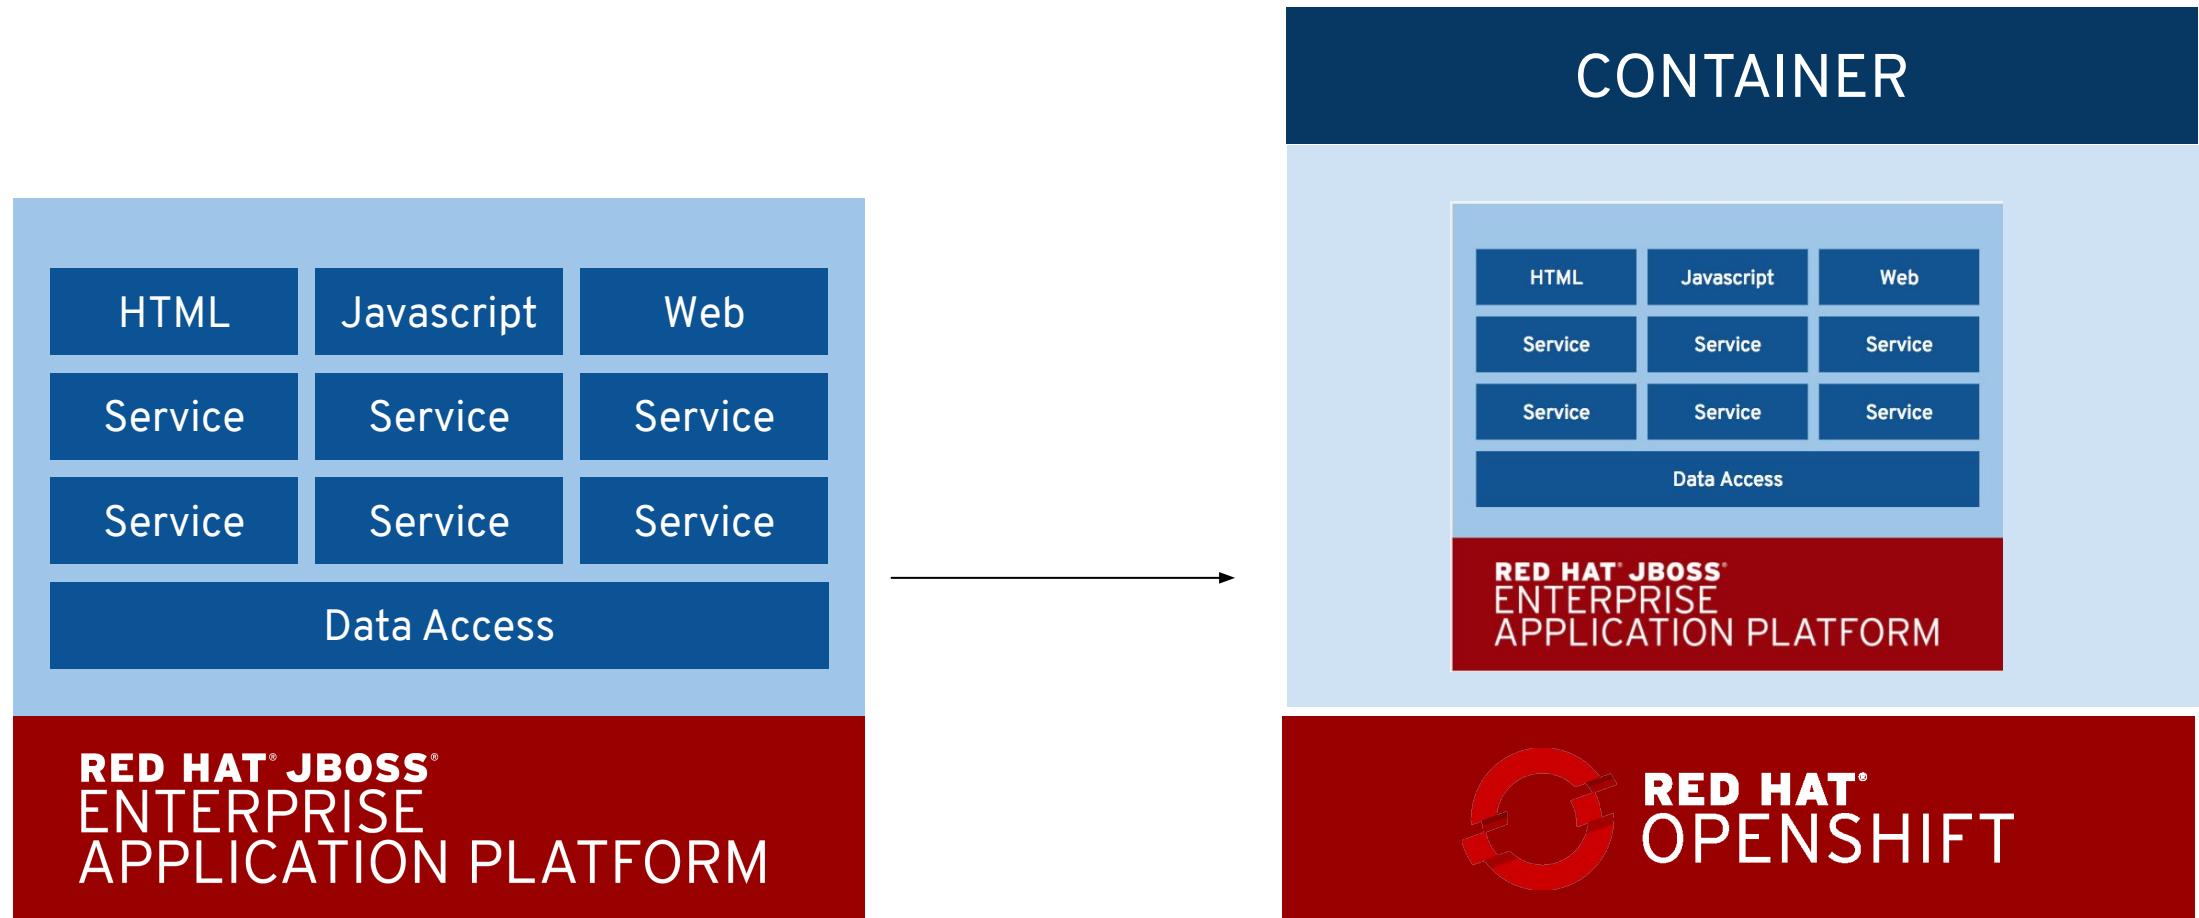
RED HAT® APPLICATION MIGRATION TOOLKIT

Catalyze large scale application modernizations and migrations

- Automate analysis
- Support effort estimation
- Accelerate code migration
- Free & Open Source



LIFT-AND-SHIFT MONOLITH TO CLOUD

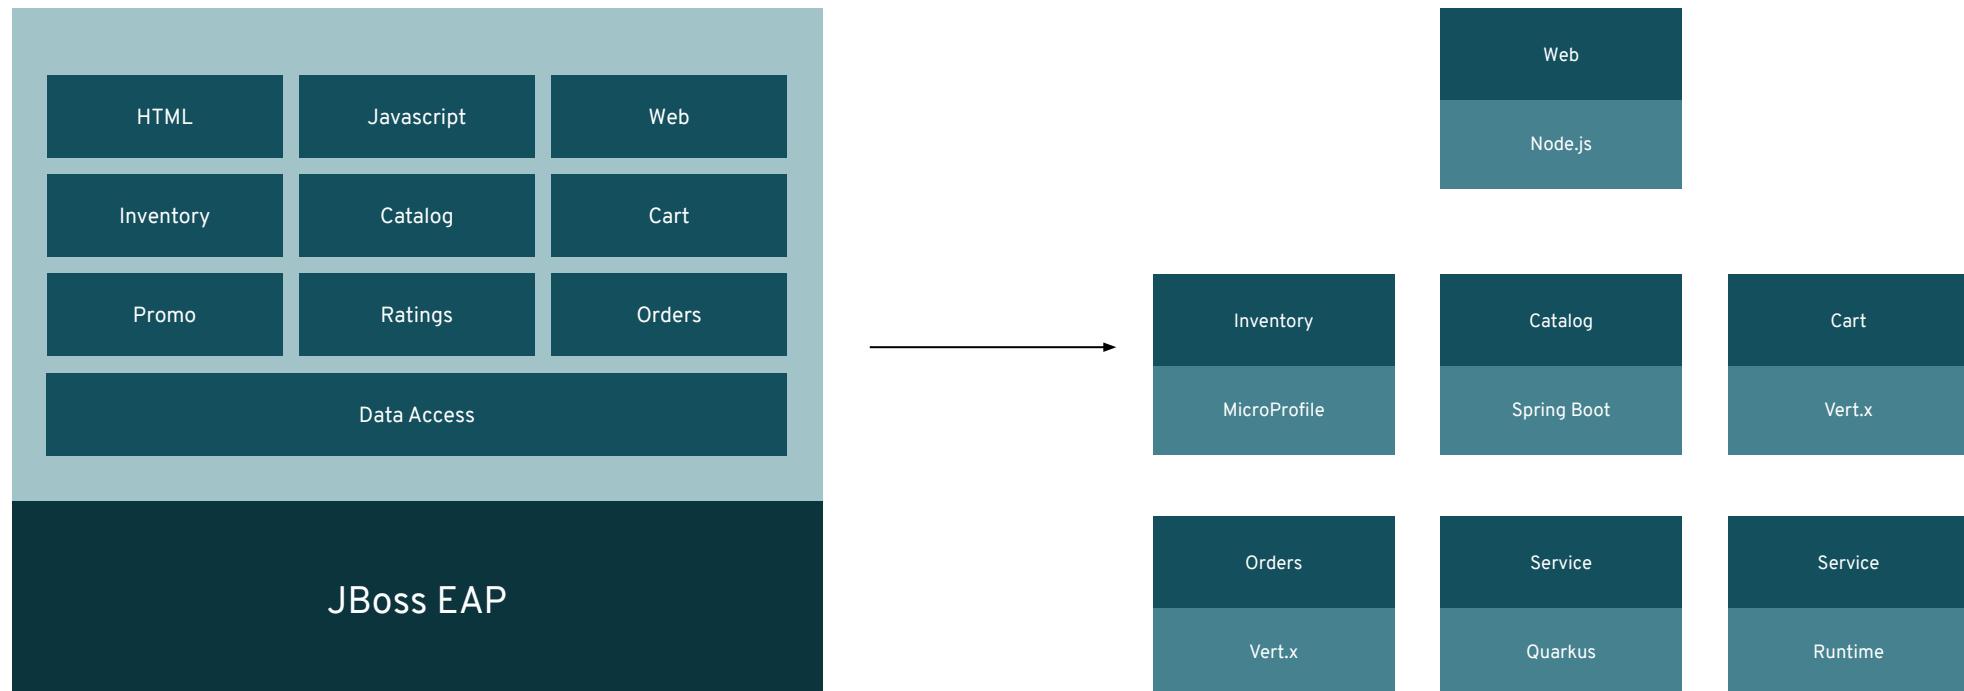


FAST-MOVING MONOLITHS

- Large organizations have a tremendous amount of resources invested in existing monolith applications
- Looking for a sane way to capture the benefits of containers and orchestration without having to complete rewrite
- OpenShift provides the platform for their existing investment with the benefit of a path forward for microservice based apps in the future

STRANGLING THE MONOLITH

- In this lab, you will begin to ‘strangle’ the coolstore monolith by implementing its services as external microservices, split along business boundaries
- As functionality is replaced, “dead” parts of monolith can be removed/retired.



CLOUD-NATIVE RUNTIMES

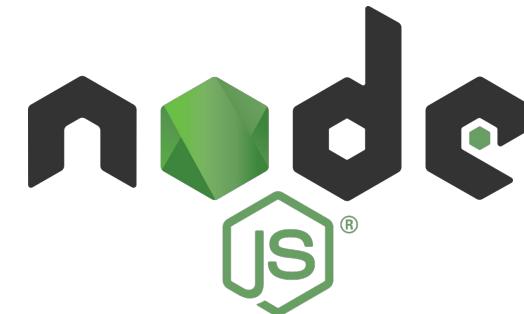


QUARKUS



spring

VERT.X





“MODERN” JAVA APP STACK

MONOLITHS, MICROSERVICES

SINGLE APP

DAYS OF LIFE

100s MB RAM

SECONDS TO START

App

Dynamic Application Frameworks

Application Server

Java Virtual Machine (Hotspot)



QUARKUS

QUARKUS

MICROSERVICES, **SERVERLESS**

SINGLE APP

SECONDS TO DAYS OF LIFE

10s MB RAM

MILLISECONDS TO START

App

Optimized Application Frameworks

Java Virtual Machine (Hotspot)



Optional





- Microservices for Developers using Spring Framework
 - Spring Boot (2.1.6), Spring Core, Spring Data, Spring Web, Spring Security, etc
- An opinionated approach to building Spring applications
- Red Hat Certified with
 - OpenShift Java Runtime
 - JBoss Web Server (Tomcat) embedded web container
- **Can also use Spring APIs in Quarkus**

GOAL FOR LAB

In this lab you will learn:

- How to use lab environment for today
- How to migrate an existing legacy Java EE application (CoolStore) from Weblogic to JBoss EAP using **Red Hat Application Migration Toolkit**
- How to deploy the result to **OpenShift container platform** to create a *Fast Moving Monolith*
- Begin modernization journey by breaking subset of monolith into microservices using **Spring Boot** and **Quarkus**

LAB INSTRUCTIONS

- **Everything is done in browser** - no local commands or installs needed on your laptop
- Tested with **Chrome 75.0.3770.142, Firefox 60.8.0esr**. → **Safari 12.x does not work!**
- If things get weird, just reload browser page
- **Turn off VPN** (we use websockets extensively), **pause AdBlock** for the lab domain (there are no ads)
- To recreate the lab locally, visit
github.com/RedHat-Middleware-Workshops/cloud-native-workshop-v2-infra
- Everyone should have their own **unique logins**, e.g.: user45 / r3dh4t1!

<http://bit.ly/spectrum-dec2019-lab1>

Credentials: userXX / r3dh4t1!

If you get stuck, raise a hand!

Module 2: Debugging, Monitoring and Continuous Delivery

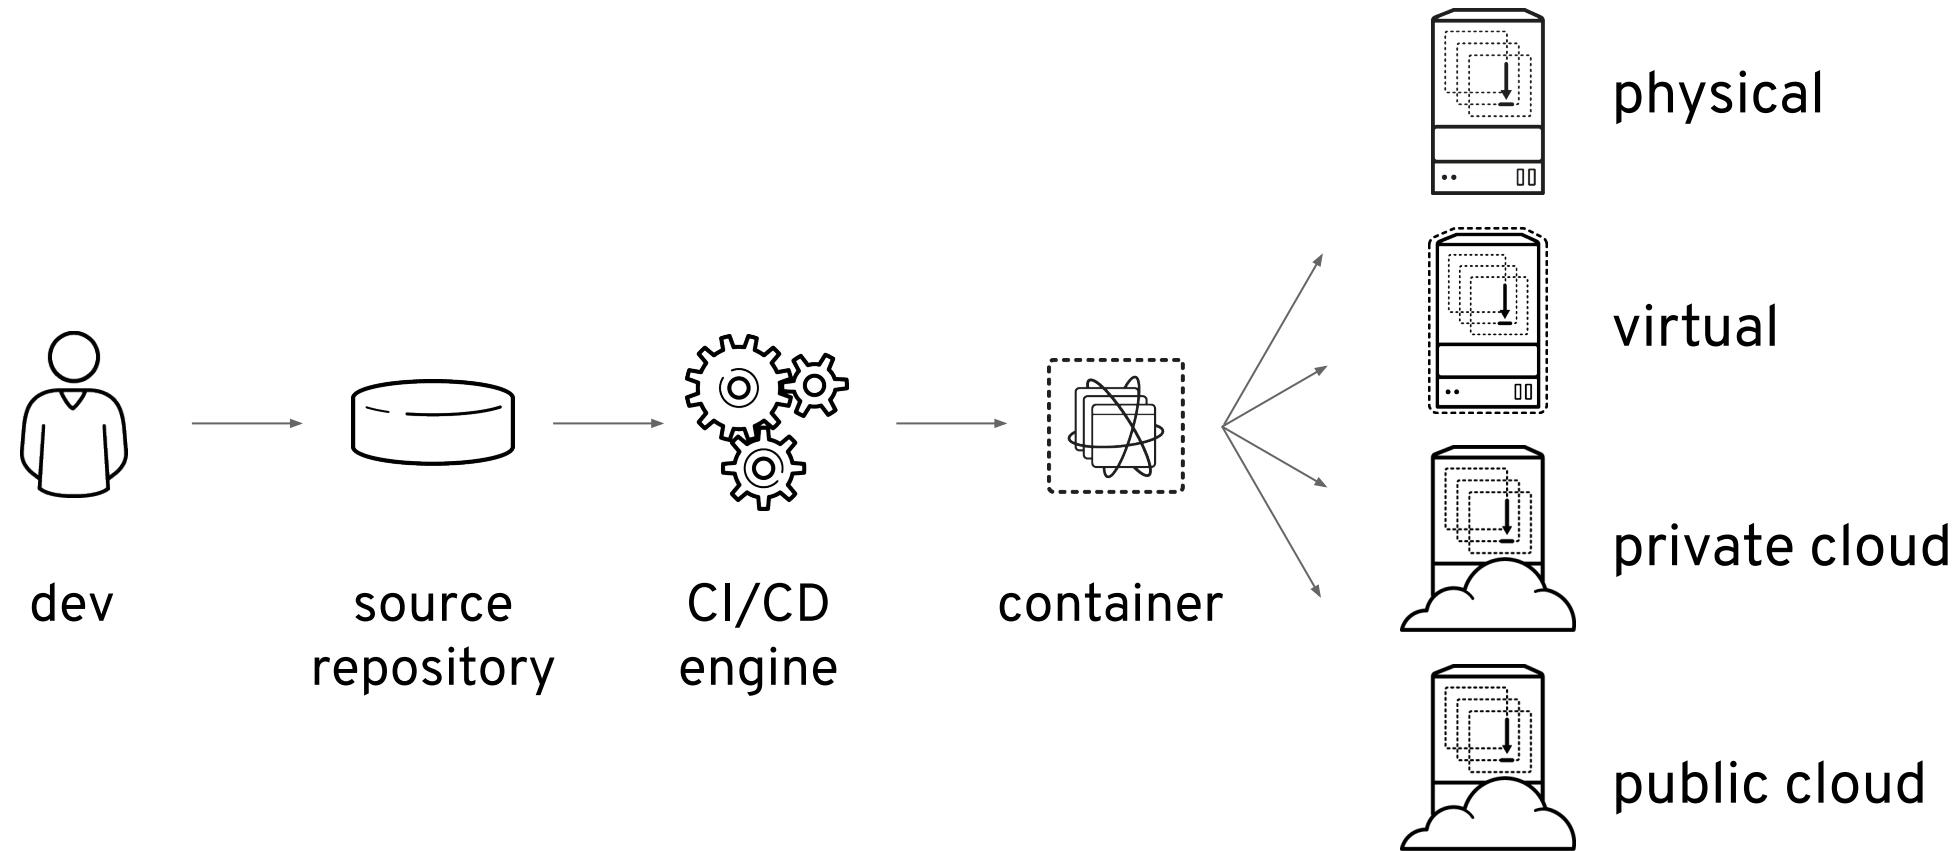
- 1. Optimizing Existing Applications
- 2. Debugging, Monitoring, and Continuous Delivery
- 3. Control Cloud Native Apps with Service Mesh
- 4. Advanced Cloud Native with Event-Driven Serverless

THE PATH TO CLOUD-NATIVE APPS

A DIGITAL DARWINISM



DEPLOYMENT PIPELINES



OPENSHIFT PIPELINES

Jenkins is still the most used CI/CD platform in enterprises and can be used from inside OpenShift.

An intuitive pipeline visualization makes it simple for users to see how builds are progressing.

The full Jenkins UI is also available.

The screenshot shows the OpenShift Pipeline interface for a Jenkins pipeline named "tasks-pipeline-1". The left sidebar has categories: Workloads, Networking, Storage, Builds (selected), Build Configs, Image Streams, Monitoring, Compute, and Administration. The main area shows the "Build Details" for "tasks-pipeline-1". The navigation tabs are Overview (selected), YAML, Environment, Logs, and Events. The "Build Overview" section shows a pipeline with the following steps and their statuses:

- Build App: Green checkmark, 6 minutes ago
- Test: Green checkmark, 5 minutes ago
- Code Analysis: Green checkmark, 5 minutes ago
- Archive App: Green checkmark, 5 minutes ago
- Build Image: Green checkmark, 5 minutes ago
- Deploy DEV: Green checkmark, 2 minutes ago

A "Promote to ST..." button is shown with an "Input Required" message and was updated 2 minutes ago. The pipeline details table includes:

NAME	STATUS
tasks-pipeline-1	Running
NAMESPACE	TYPE
NS cicd-smx	JenkinsPipeline

OPENShift PIPELINES

CONFIDENTIAL Designator

- CI/CD workflow via Jenkins
- Pipelines are started, monitored, and managed similar to other builds
- Auto-provisioning of Jenkins server
- On-demand Jenkins slaves
- Embedded Jenkinsfile or in Git repo

```
pipeline {
    agent {
        label 'maven'
    }
    stages {
        stage('build app') {
            steps {
                git url:
'https://git/app.git'
                sh "mvn package"
            }
        }
        stage('build image') {
            steps {
                script {
                    openshift.withCluster() {
                        openshift.startBuild("...")
                    }
                }
            }
        }
    }
}
```

OPENSHIFT PIPELINES CI/CD PLATFORM

CONFIDENTIAL Designator

Next-gen Kubernetes CI/CD pipeline that works for containers (including serverless).

Based on the **Tekton** project (which was spun out of the Knative Pipelines project) started by Google, Red Hat and others.

The screenshot shows the Red Hat OpenShift Pipelines interface. At the top, there's a navigation bar with the Red Hat logo and 'OPENSHIFT'. On the right, it shows 'Administrator' and a user icon. Below the navigation, there's a search bar with 'XYZ Name' and a dropdown for 'Project: Default'. A filter bar at the top right shows '2 All' items, with buttons for Pending, Running, Complete, Failed, Error, and Cancelled, along with 'Select All Filters' and a count of '2 Items'. The main area is titled 'Builds' and shows a single pipeline named 'aa-build-3' which is 'Running' and started '10 mins ago' with a duration of '2 min 04 sec'. The pipeline has a commit trigger labeled 'Commit #123456ABC'. Below the pipeline name, there's a detailed view of the pipeline steps:

```
graph LR; input((input info)) --> buildName[build-name (30s)]; buildName --> test1[Test-st... (6s)]; test1 --> codeA[Code a... (13s)]; codeA --> imageB[Image b... (0s)]; imageB --> deploy[DeployTo... (0s)];
```

The steps are color-coded: green for build-name, blue for Test-st..., blue for Code a..., grey for Image b..., and grey for DeployTo... . Step details are as follows:

- build-name (30s)**: Status: Success, Duration: 30s, Substeps: 3/3
- Test-st... (6s)**: Status: Success, Duration: 6s, Substeps: 3/8
- Code a... (13s)**: Status: Success, Duration: 13s, Substeps: 2/5
 - Steps - (1/3)**: Status: Success, Duration: 1s
 - Security... (20s)**: Status: Success, Duration: 20s
- Image b... (0s)**: Status: Success, Duration: 0s, Substeps: 0/2
- DeployTo... (0s)**: Status: Success, Duration: 0s, Substeps: 0/7

At the bottom of the pipeline view, there's a terminal window showing the command-line output of the pipeline execution:

```
Downloading six-1.11.0-py2.py3-none-any.whl
Building wheels for collected packages: tornado, configparser
Running setup.py bdist_wheel for tornado: started
Running setup.py bdist_wheel for tornado: finished with status 'done'
Stored in directory: /root/.cache/pip/wheels/0c/21/02/8cd6a381450df92b449ea7c57be653dd7aa80ba42c716212c
Running setup.py bdist_wheel for configparser: started
Running setup.py bdist_wheel for configparser: finished with status 'done'
Stored in directory: /root/.cache/pip/wheels/1c/bd/b4/277af3f6c40645661b4cd1c21df26aca0f2e1e9714a1d4cda8
Successfully built tornado configparser
Installing collected packages: six, singledispatch, certifi, backports-abc, tornado, enum34, configparser, mccabe, pyflakes, pycodestyle, flake8
  Found existing installation: six 1.8.0
    Uninstalling six-1.8.0:
      Successfully uninstalled six-1.8.0
Successfully installed backports-abc-0.5 certifi-2017.11.5 configparser-3.5.0 enum34-1.1.6 flake8-3.5.0 mccabe-0.6.1 pycodestyle-2.3.1 pyflakes-1.6.0
singledispatch-3.4.0.3 six-1.11.0 tornado-4.5.3
$ python -c 'print("Hello, world")'
Hello, world
Job succeeded
```

OpenShift Application Monitoring



Metrics collection and storage via Prometheus, an open-source monitoring system time series database.

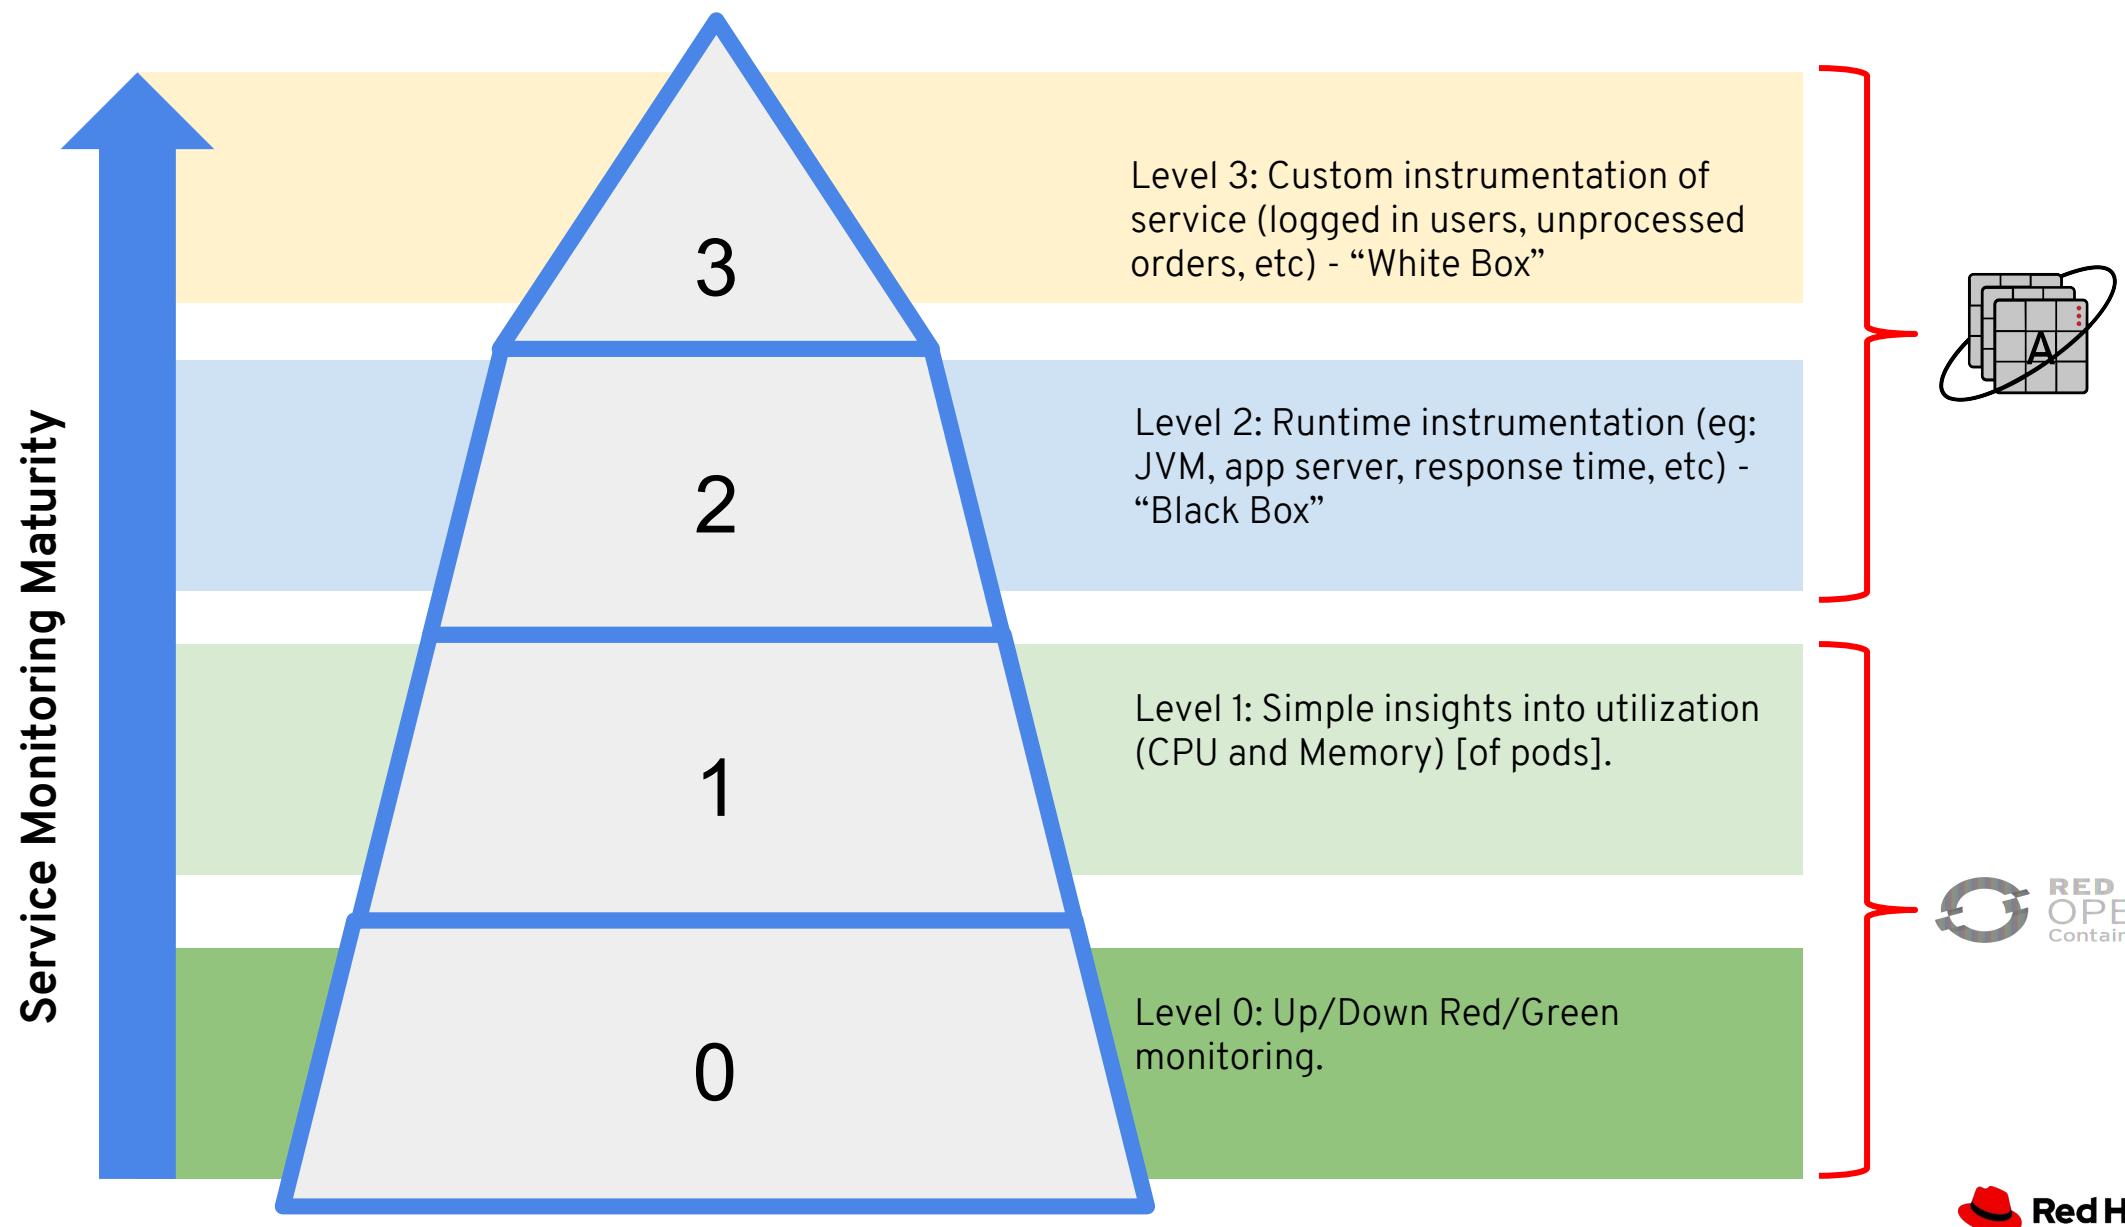


Alerting/notification via Prometheus' Alertmanager, an open-source tool that handles alerts sent by Prometheus.

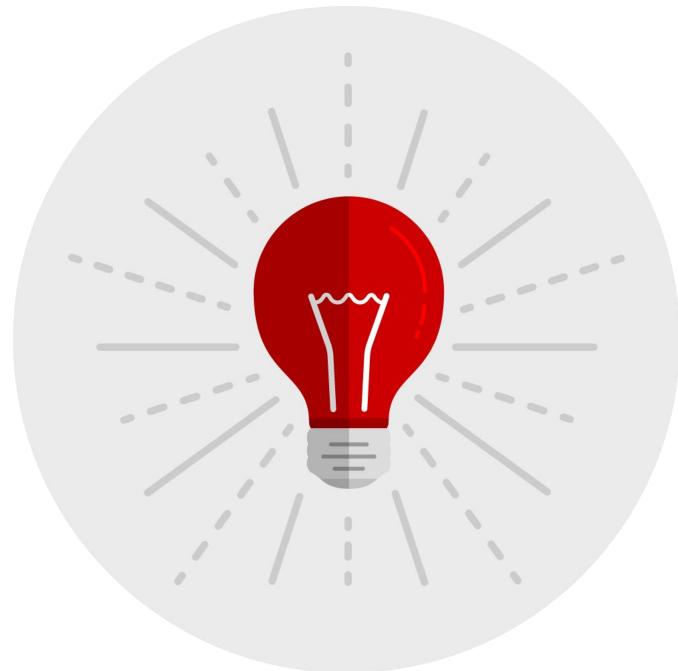


Metrics visualization via Grafana, the leading metrics visualization technology.

Service Monitoring Maturity Model



Middleware OOTB Instrumentation



Prometheus Endpoint

Runtimes (Thorntail, Node.js, vert.x, Quarkus, Spring Boot),
Decision & Process Servers.
Coming in EAP 7.3

Health Endpoint

EAP, Runtimes (Thorntail, Node.js, vert.x, Spring Boot,
Quarkus), Decision & Process Servers.

Distributed Tracing Client

EAP, Runtimes (Thorntail, Node.js, vert.x, Spring Boot,
Quarkus), Fuse.

JSON logging format

EAP, Spring Boot, Decision & Process Servers

GOAL FOR MODULE

In this module you will learn:

- How to create a CI/CD pipeline using Jenkins to automate the deployment of updates to microservice applications
- How to use CodeReady Workspaces to debug Java apps
- How to create health probes to allow the platform to act on unhealthy apps
- How to use distributed tracing and Jaeger to gain insight into application behavior
- How to instrument and monitor application performance using Prometheus (metrics ingestion and alerting) and Grafana (visualization)

LAB INSTRUCTIONS

- **Everything is done in browser** - no local commands or installs needed on your laptop
- Tested with **Chrome 75.0.3770.142, Firefox 60.8.0esr**. → **Safari 12.x does not work!**
- If things get weird, just reload browser page
- **Turn off VPN** (we use websockets extensively), **pause AdBlock** for the lab domain (there are no ads)
- To recreate the lab locally, visit
github.com/RedHat-Middleware-Workshops/cloud-native-workshop-v2-infra
- Everyone should have their own **unique logins**, e.g.: user45 / r3dh4t1!

Get Started at [LINK]

Credentials: userXX / r3dh4t1!

If you get stuck, raise hand

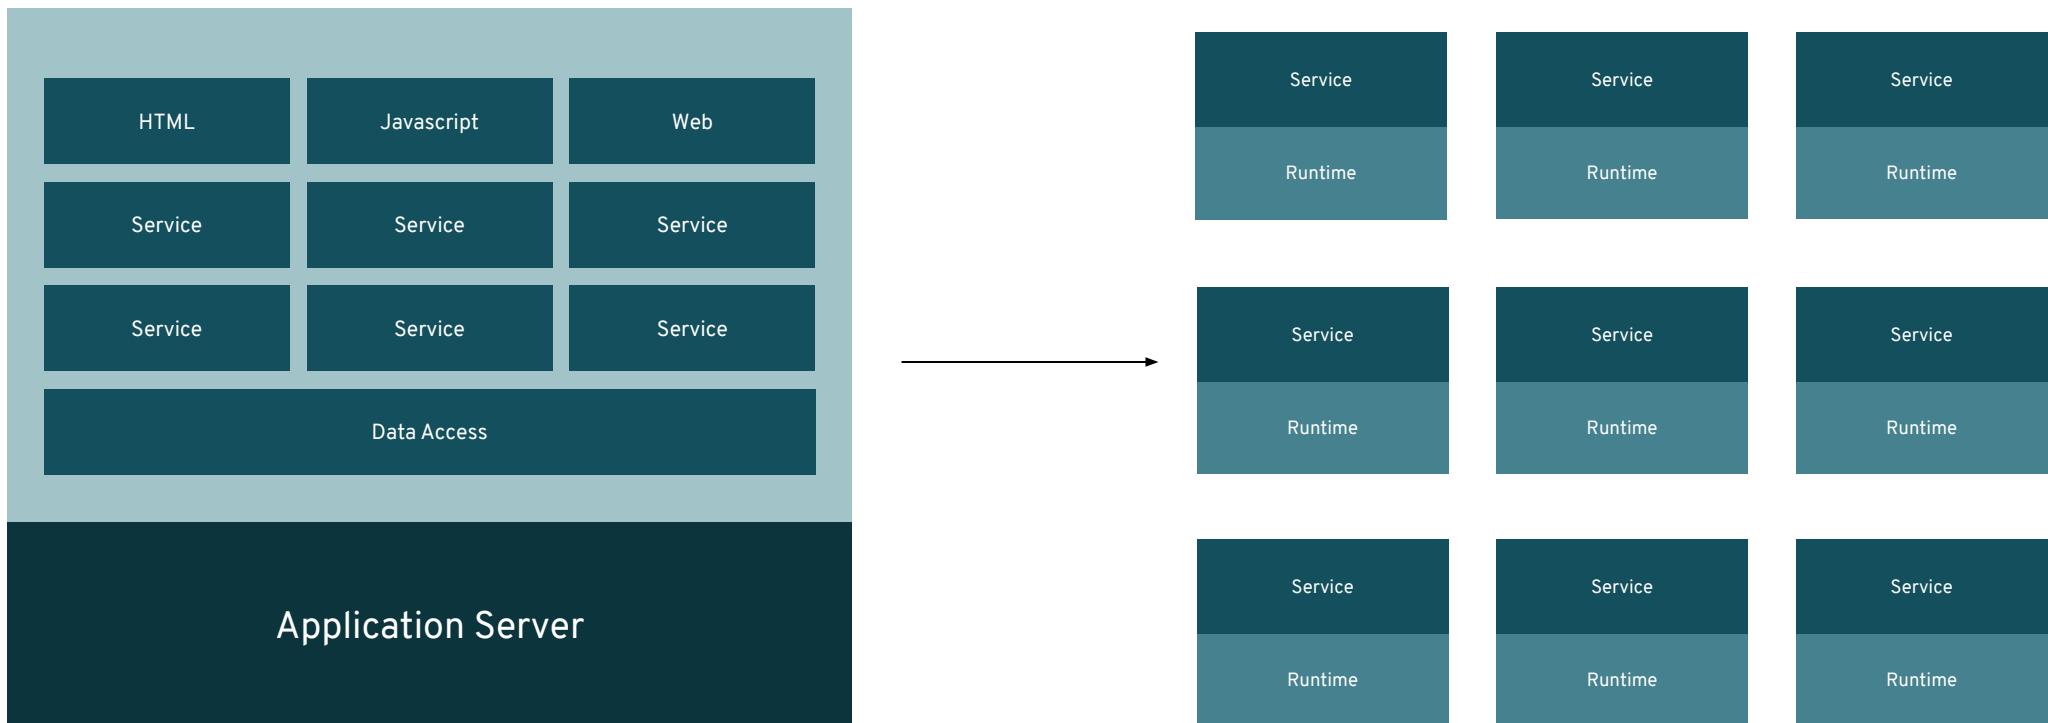
Module 3: Control Cloud Native Apps With Service Mesh

- 1. Optimizing Existing Applications
- 2. Debugging, Monitoring, and Continuous Delivery
- 3. Control Cloud Native Apps with Service Mesh
- 4. Advanced Cloud Native with Event-Driven Serverless



MICROSERVICES ARCHITECTURE DISTRIBUTED

CONFIDENTIAL Designator



DISTRIBUTED COMPUTING CHALLENGES

Fallacies of Distributed Computing

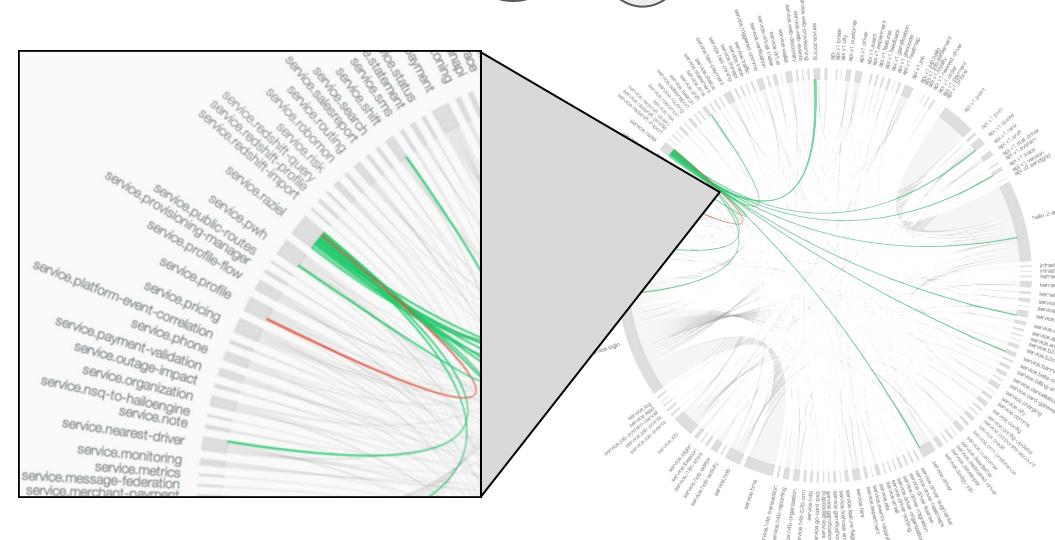
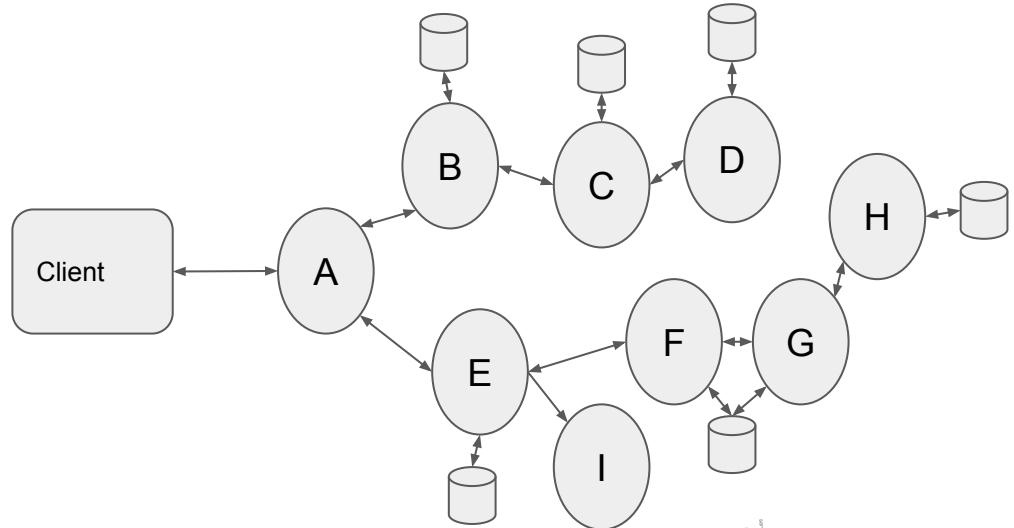
- The network is reliable.
- Latency is zero.
- Bandwidth is infinite.
- The network is secure.
- Topology doesn't change.
- There is one administrator.
- Transport cost is zero.
- The network is homogeneous.

[wikipedia.org/wiki/Fallacies_of_distributed_computing](https://en.wikipedia.org/wiki/Fallacies_of_distributed_computing)

MICROSERVICES ARE HARD

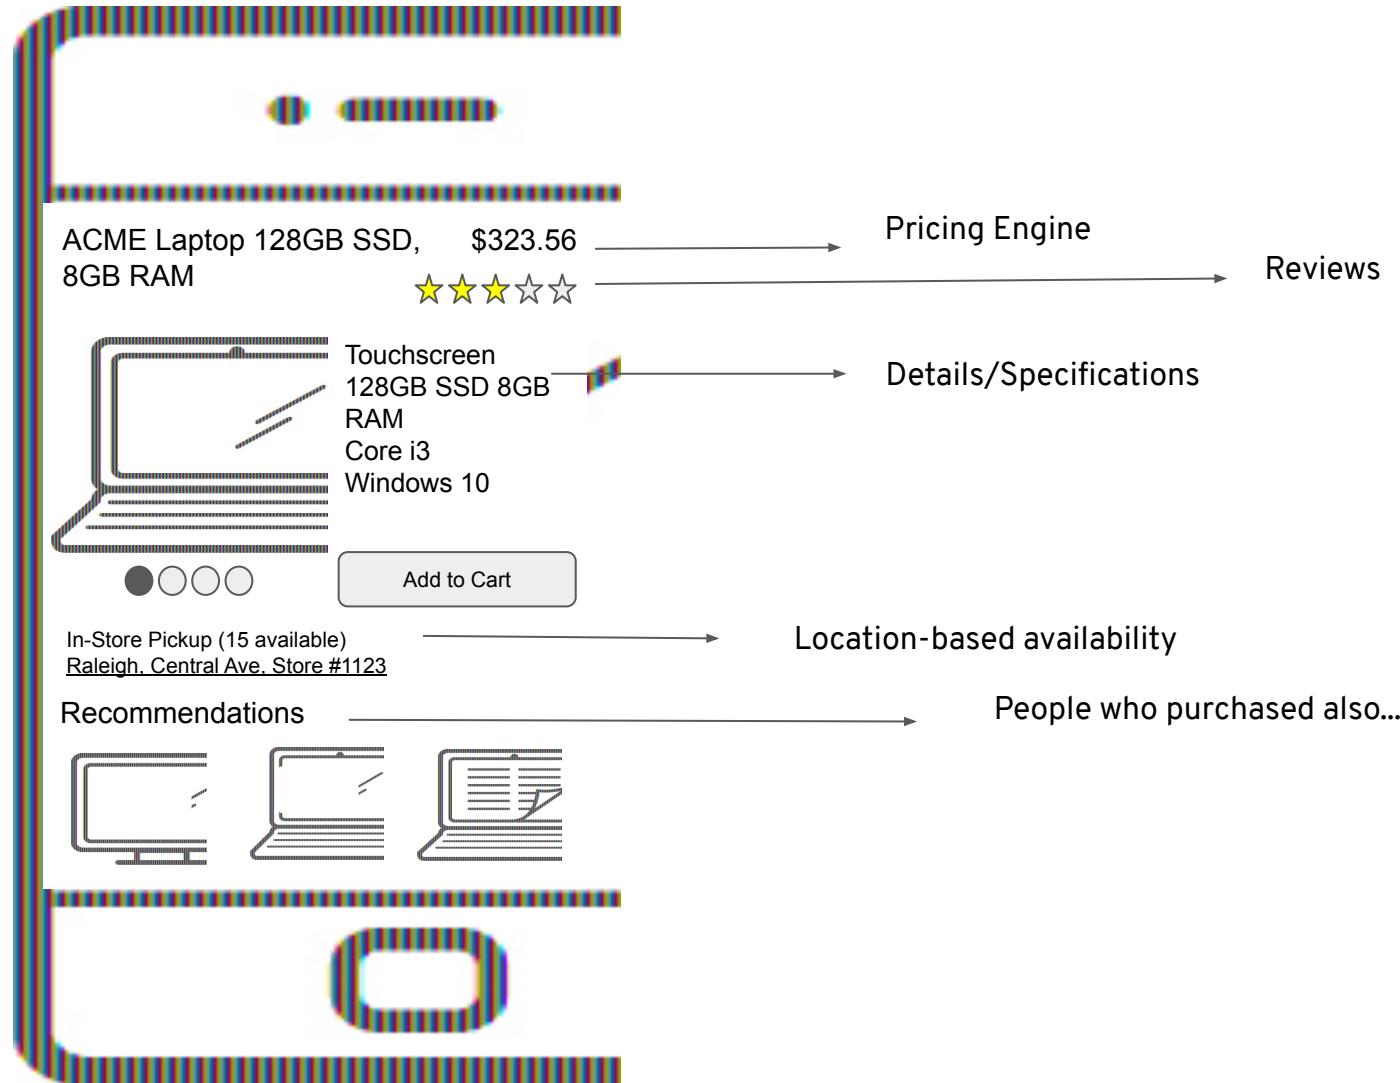
Because applications must deal with

- Unpredictable failures
- End-to-end application correctness
- System degradation
- Topology changes
- Elastic/ephemeral/transient resources
- Distributed logs
- The fallacies of distributed computing



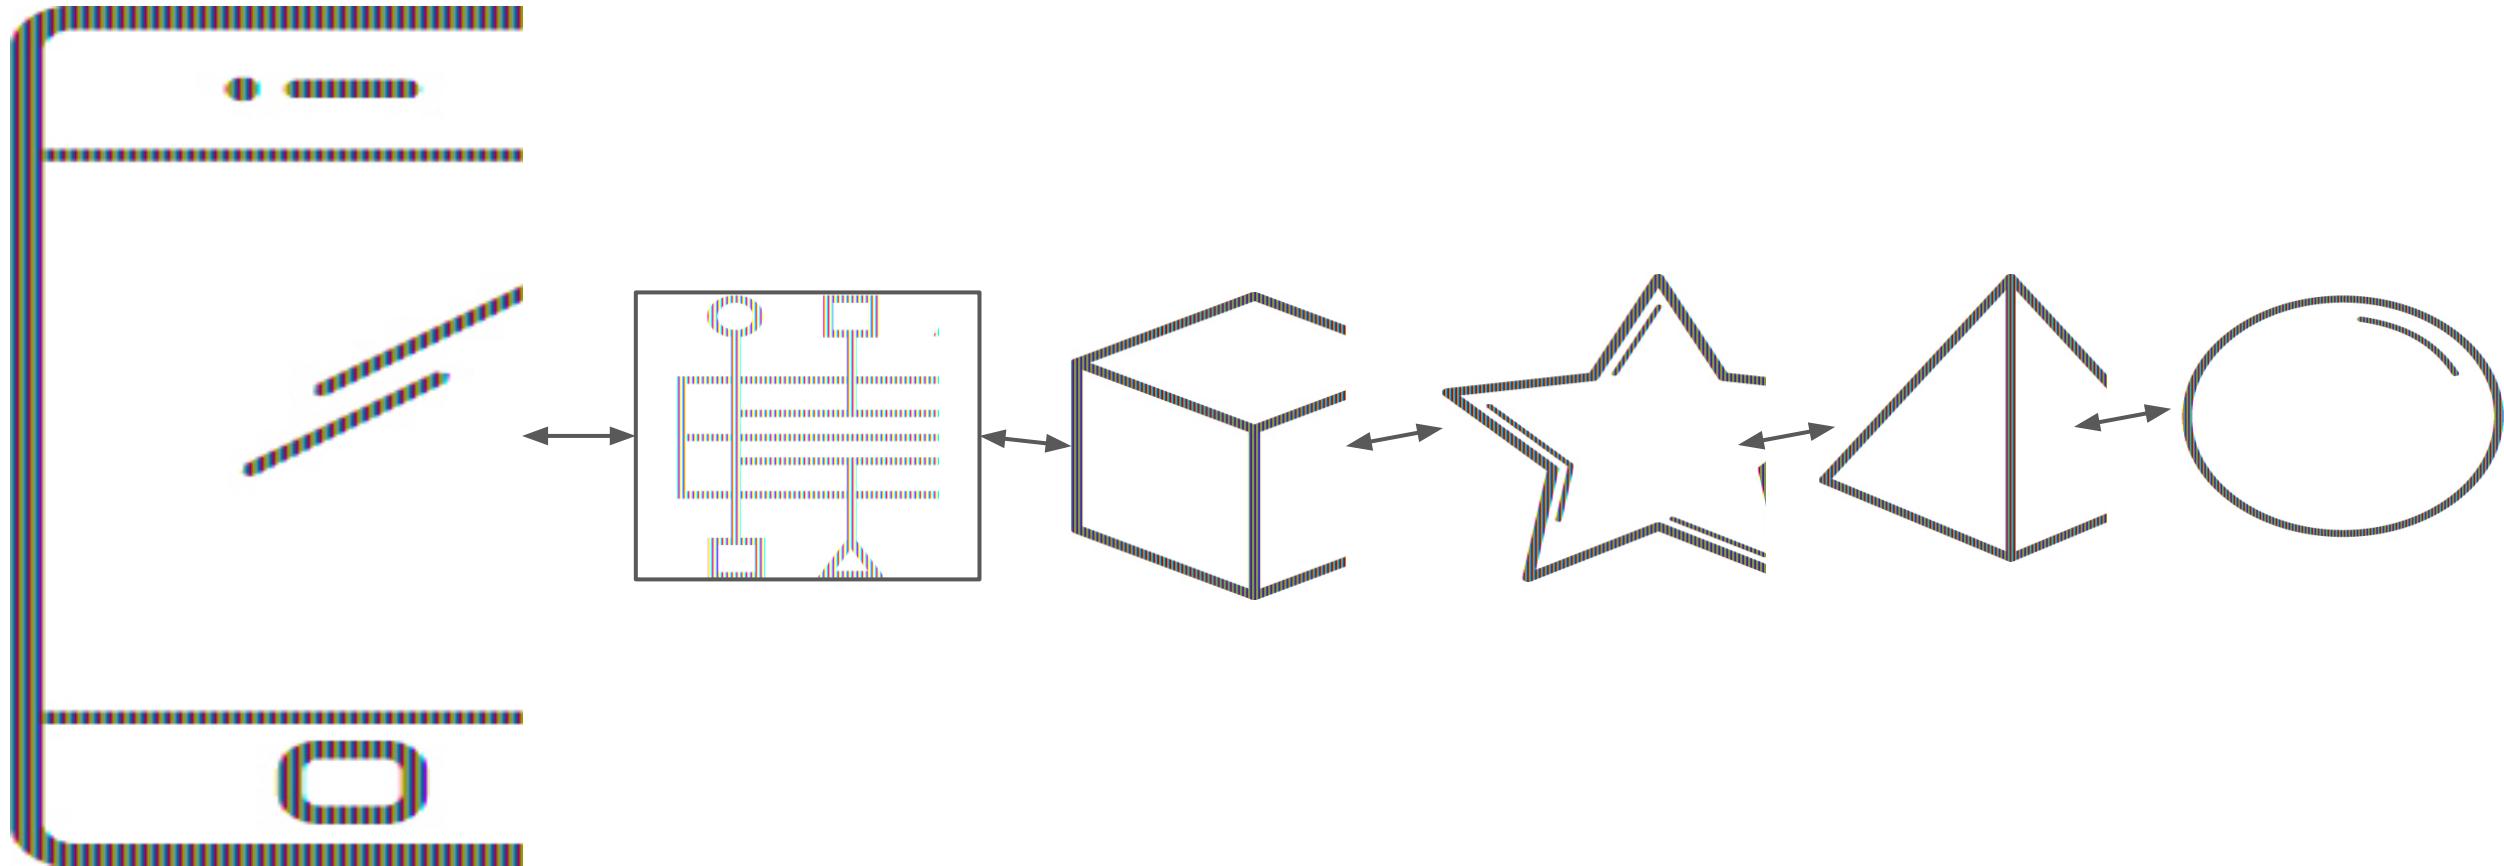
AN EXAMPLE

CONFIDENTIAL Designator

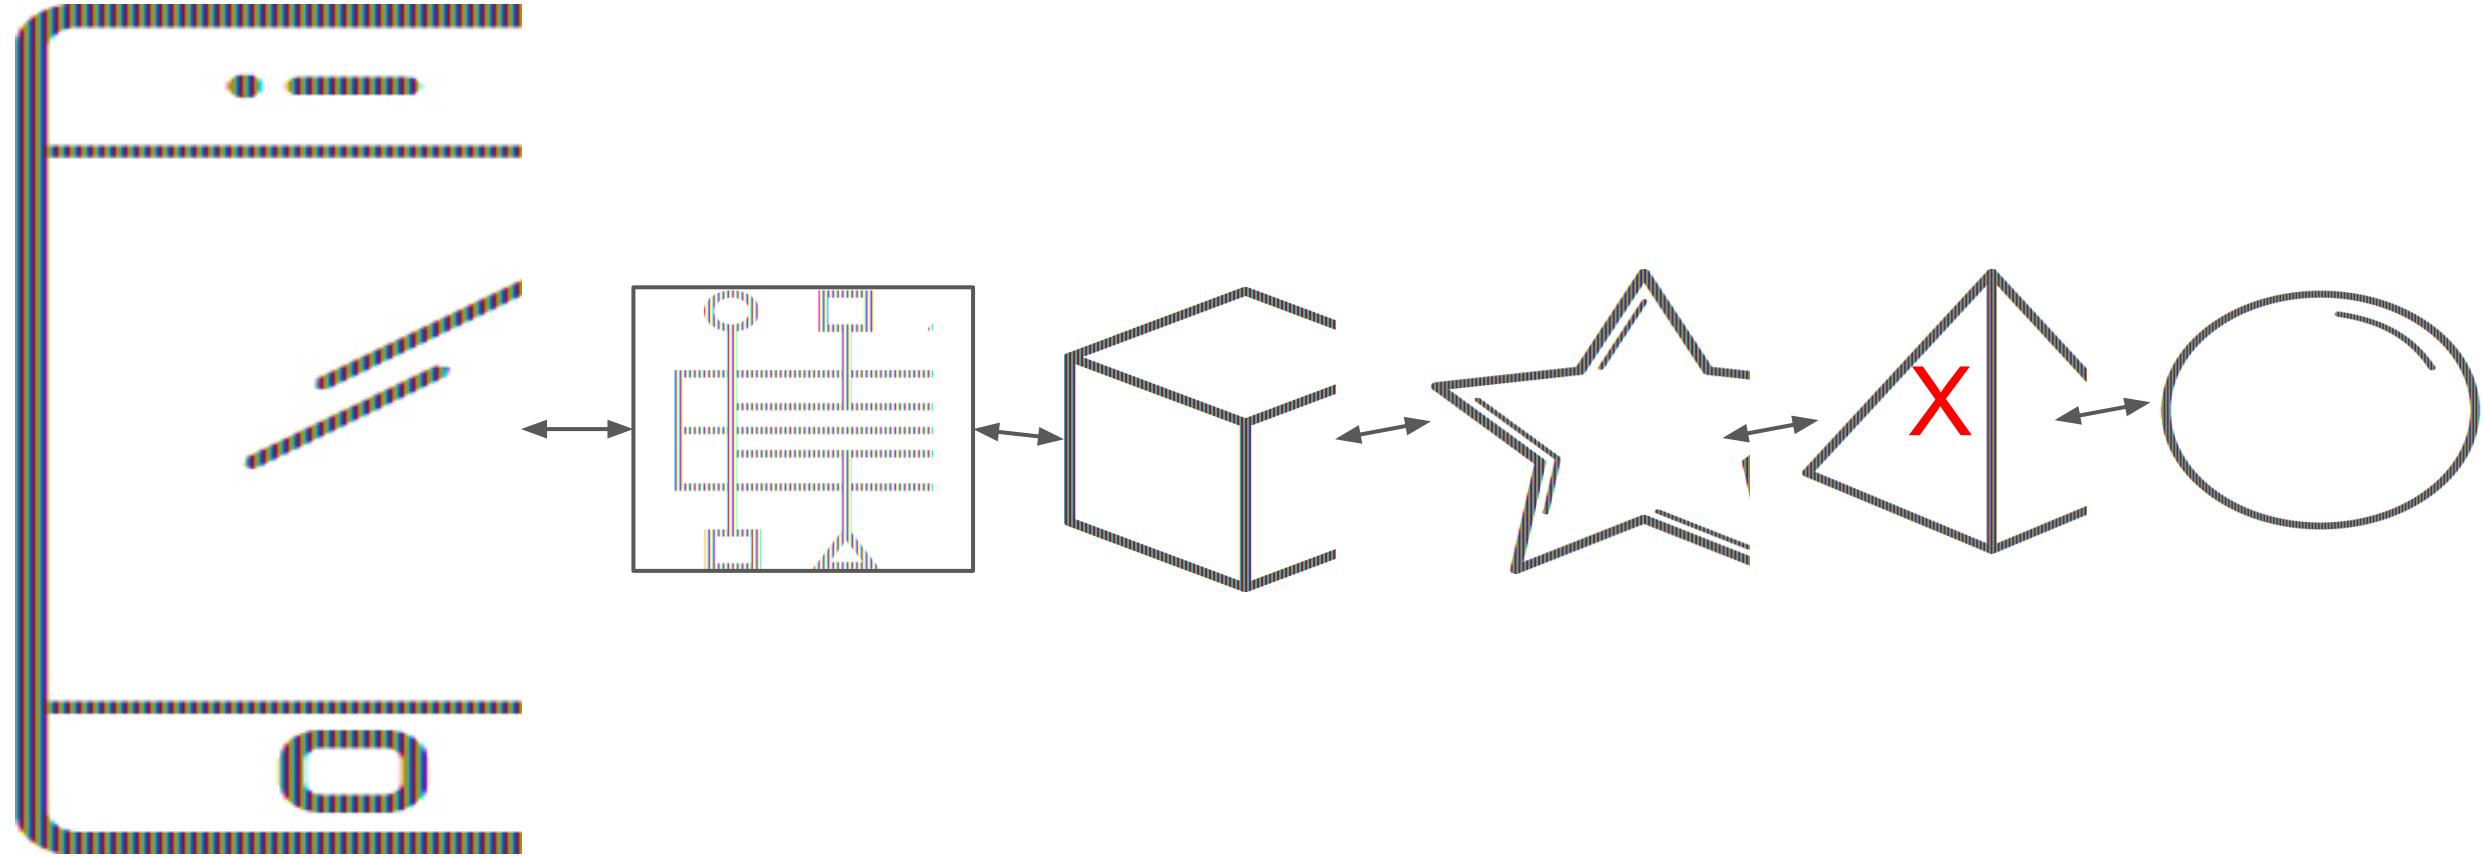


CHAINING

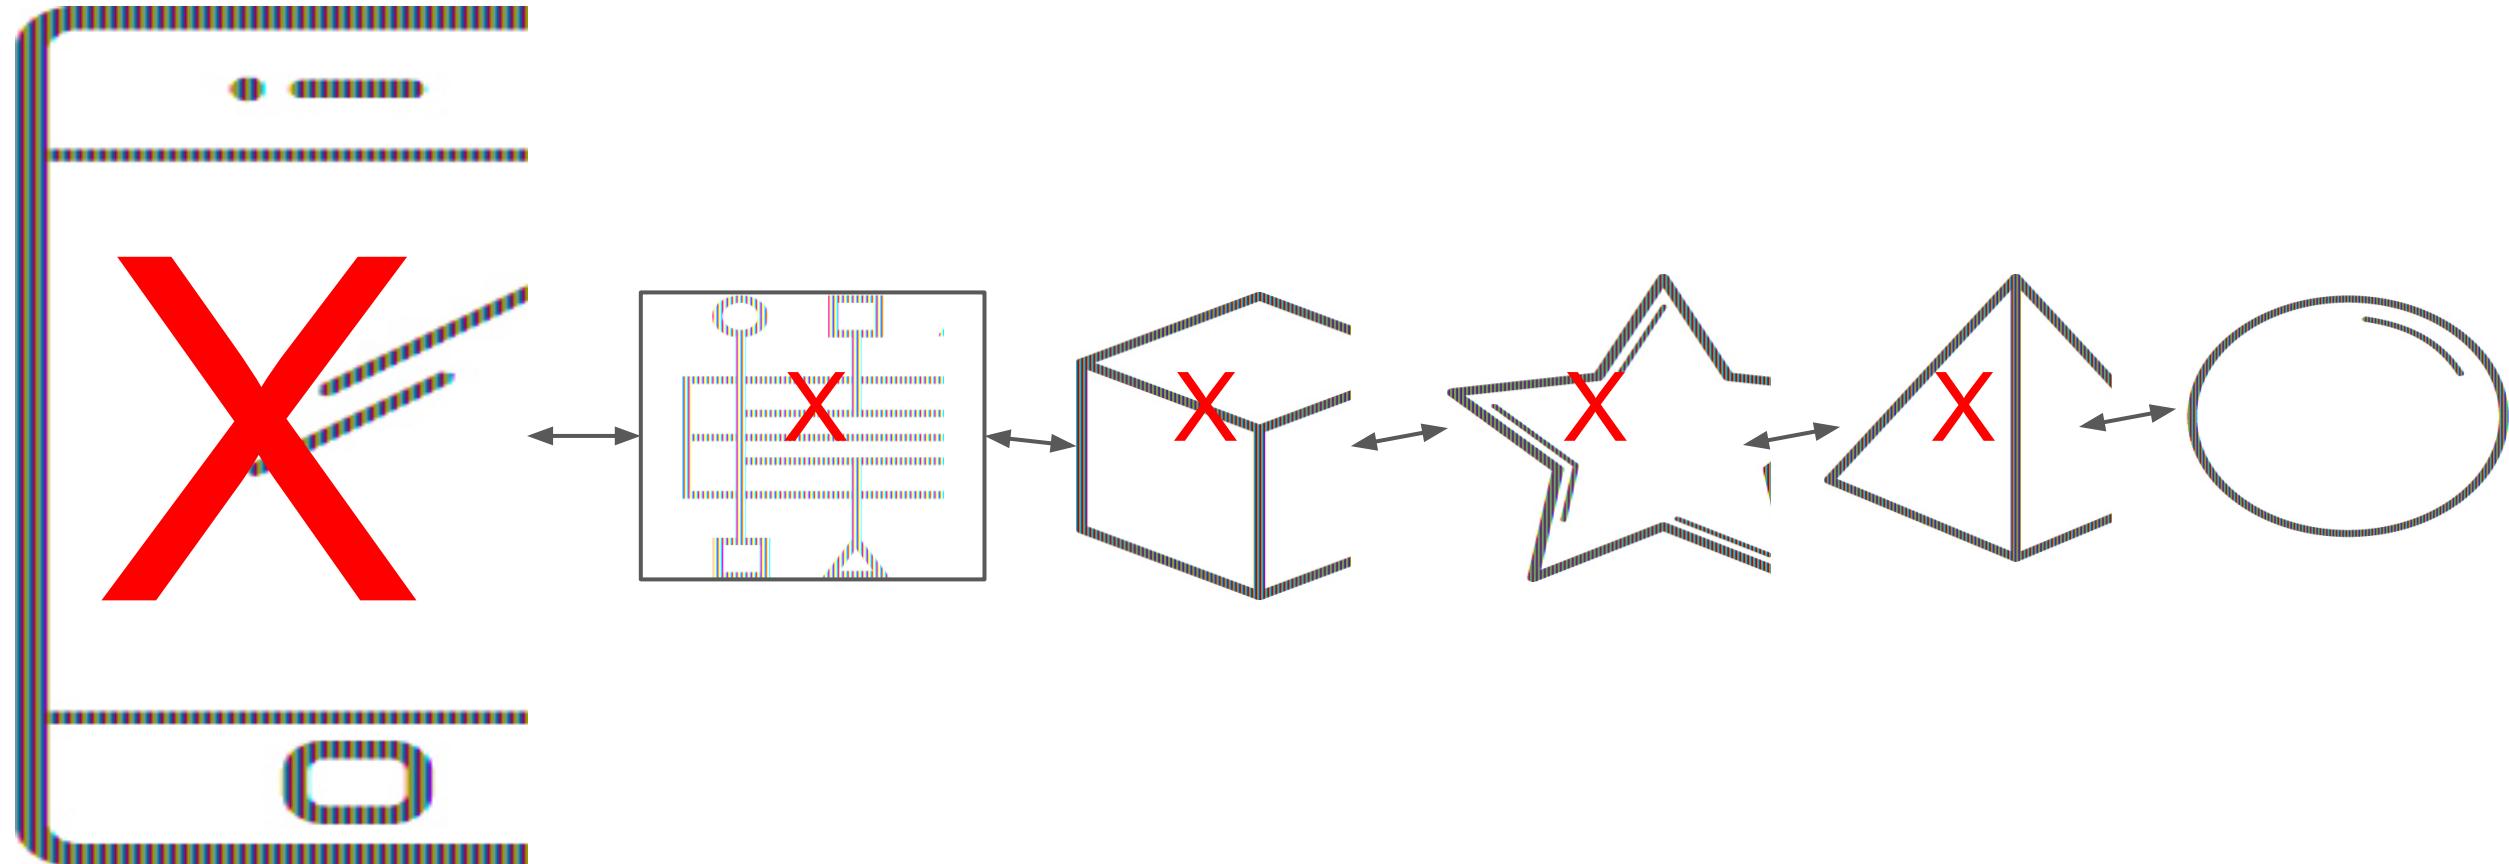
CONFIDENTIAL Designator



CHAINING (FAILURE)



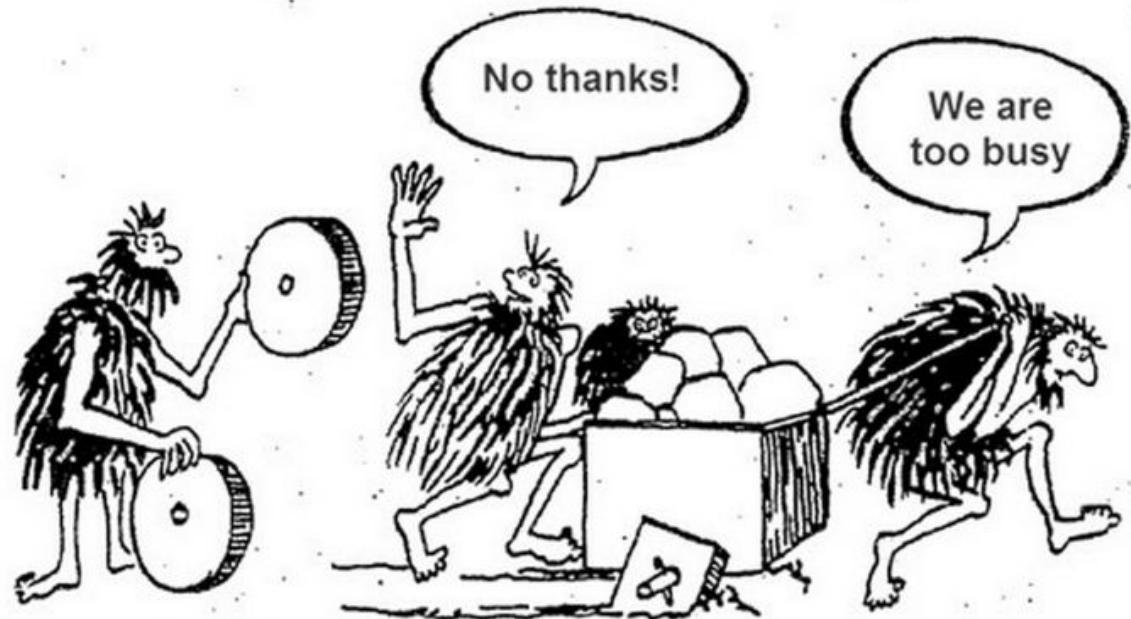
CHAINING (CASCADING FAILURE)

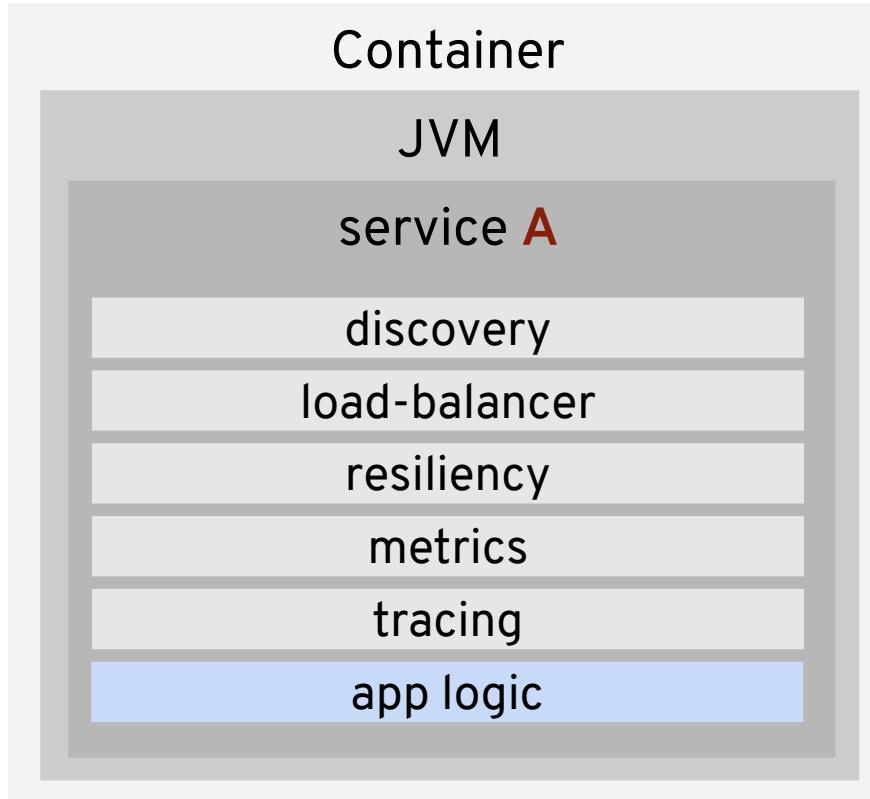


POSSIBLE SOLUTIONS

Have your developers do this:

- Circuit Breaking
- Bulkheading
- Timeouts/Retries
- Service Discovery
- Load Balancing
- Traffic Control

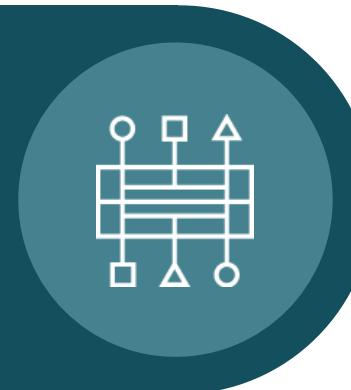




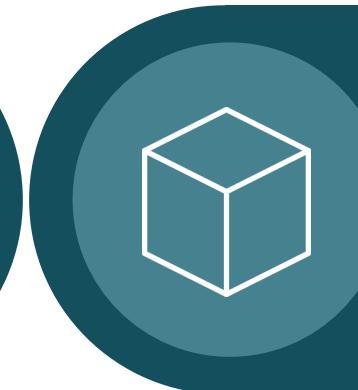
Need a library to support each language/framework combination

WHAT ABOUT...?

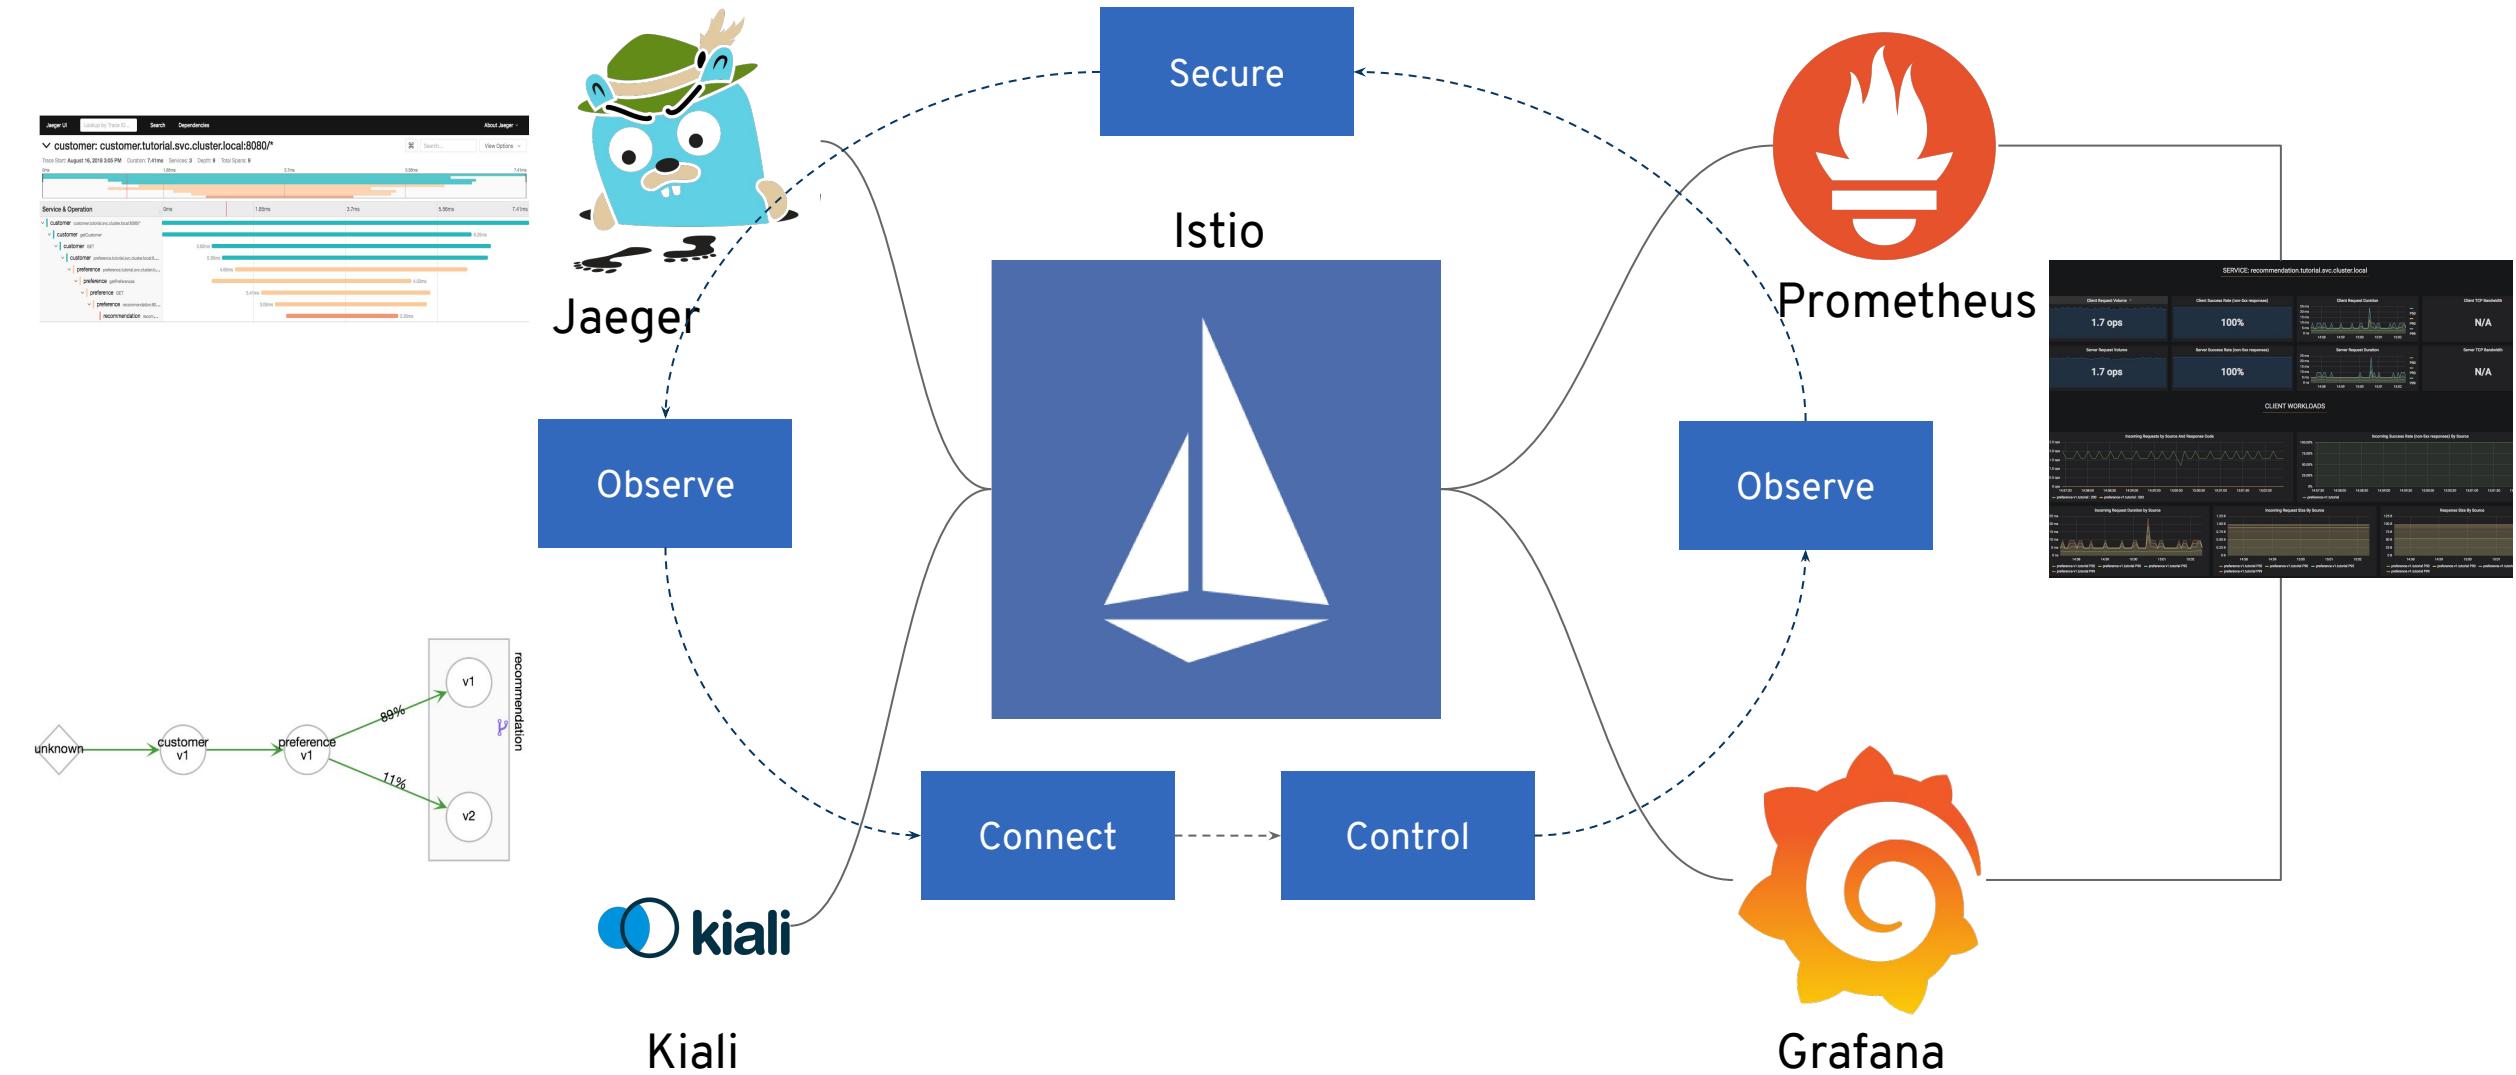
POLYGLOT
APPS



EXISTING
APPS



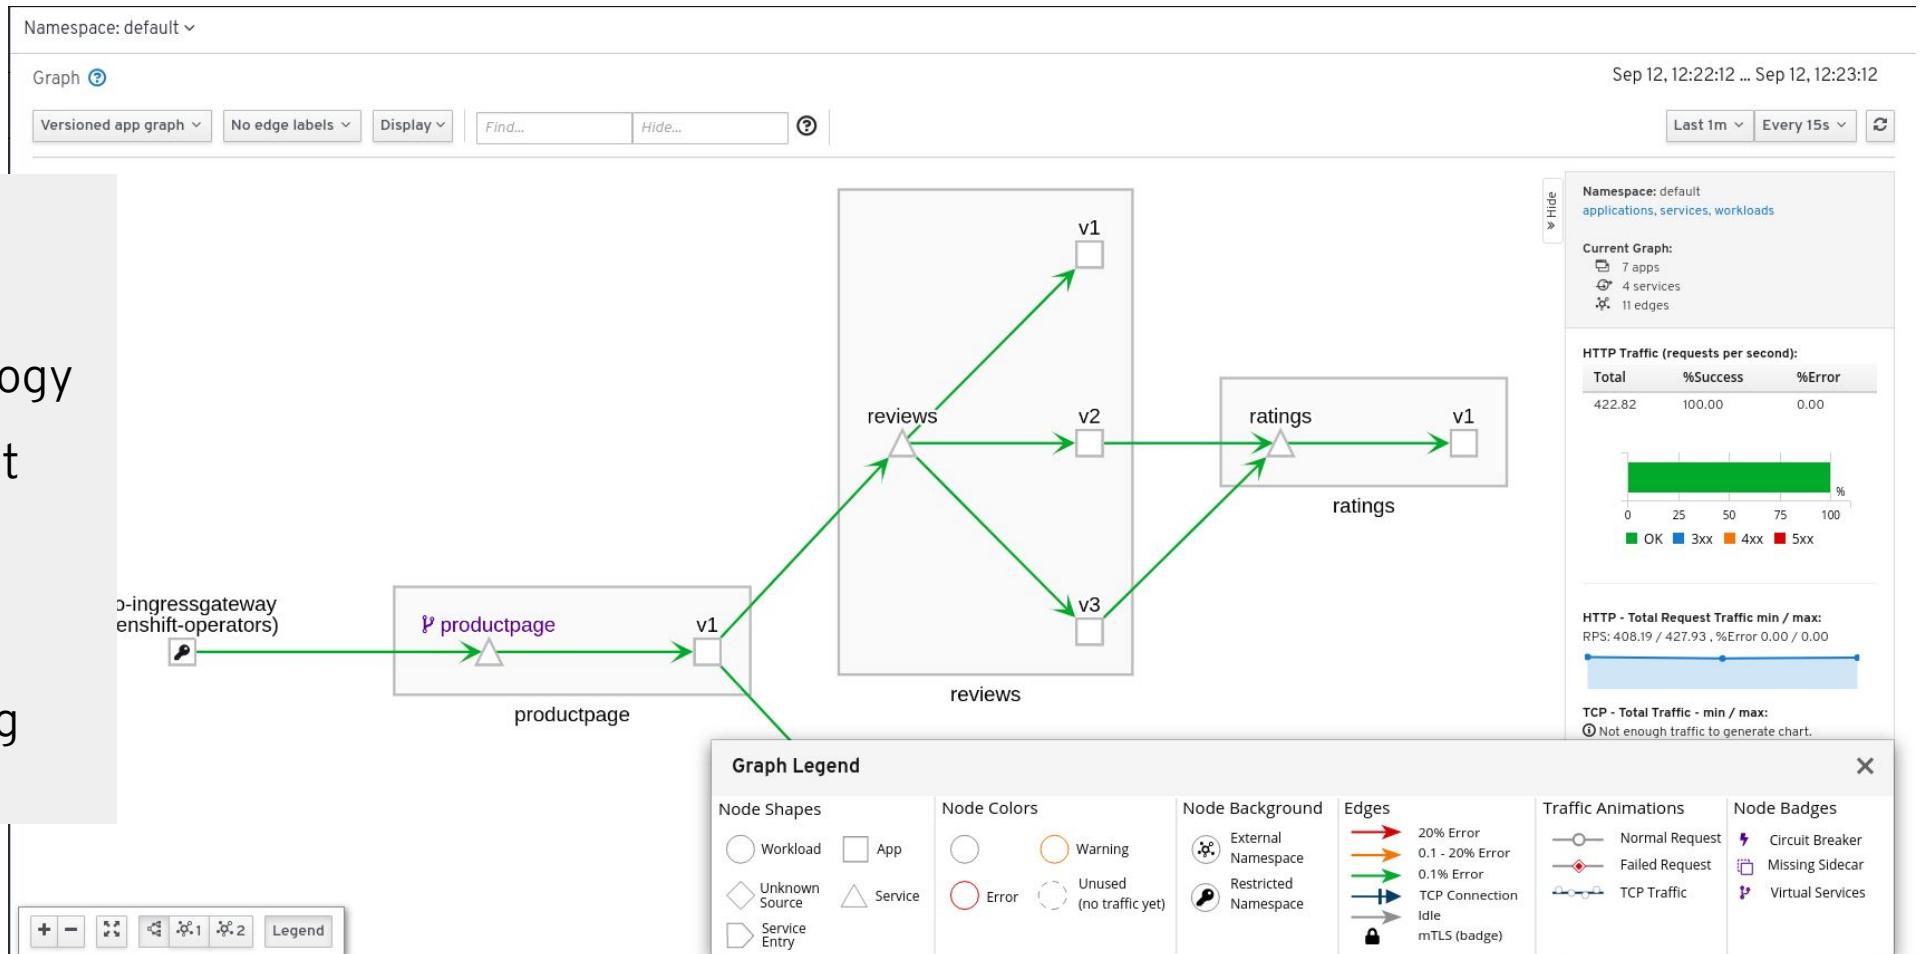
Service Mesh Ecosystem



Enhanced Visualization of Cluster Traffic With Kiali

Visualization of what
Matters most:

- Application Topology
- Traffic throughput
- Error Rates
- Service Latency
- Service Versioning



GOAL FOR LAB

In this lab you will learn:

- How to deploy apps into the **OpenShift Service Mesh**
- How to generate and visualize **deep metrics** for apps with **Kiali** console
- How to **alter routing** dynamically
- How to **inject faults** for testing
- How to do **rate limiting**
- How the mesh implements **circuit breaking** and **distributed tracing**

LAB INSTRUCTIONS

- **Everything is done in browser** - no local commands or installs needed on your laptop
- Tested with **Chrome 75.0.3770.142, Firefox 60.8.0esr**. → **Safari 12.x does not work!**
- If things get weird, just reload browser page
- **Turn off VPN** (we use websockets extensively), **pause AdBlock** for the lab domain (there are no ads)
- To recreate the lab locally, visit
github.com/RedHat-Middleware-Workshops/cloud-native-workshop-v2-infra
- Everyone should have their own **unique logins**, e.g.: user45 / r3dh4t1!

<http://bit.ly/spectrum-dec2019-lab3>

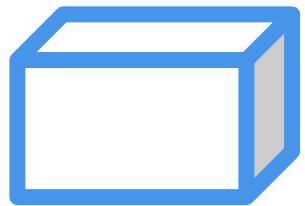
Credentials: userXX / r3dh4t1!

If you get stuck, raise a hand!

Module 4: Advanced Cloud Native Architectures With Event-driven Serverless

- 1. Optimizing Existing Applications
- 2. Debugging, Monitoring, and Continuous Delivery
- 3. Control Cloud Native Apps with Service Mesh
- 4. Advanced Cloud Native with Event-Driven Serverless

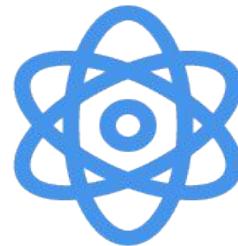
Growth in Application Architecture Choices



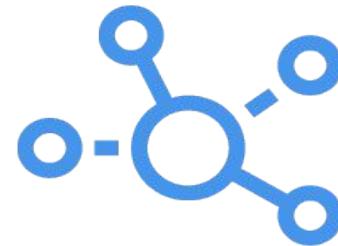
Monolith



Cloud Native



Microservices

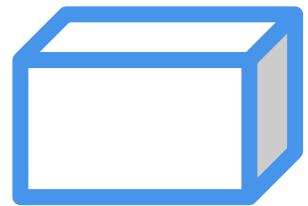


Serverless

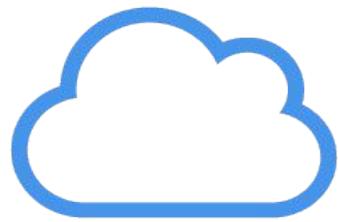


**Event-Driven
Architecture**

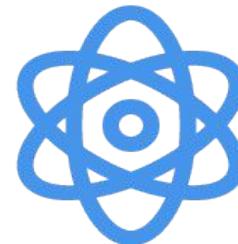
Common Deployment Platform



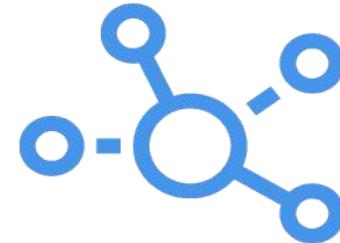
Monolith



Cloud Native



Microservices



Serverless



**Event-Driven
Architecture**



kubernetes



Istio



Knative

RISE IN EVENTS DRIVEN BY MODERN APP-DEV

CONFIDENTIAL Designator

1

MULTI-CLOUD

Deploying applications across on-premise and public cloud drives need to sync state and notify dependent applications anywhere

2

NEAR-REALTIME

End users expect a near realtime experience from modern applications

3

AVAILABILITY ISOLATION

Deployments must be resilient by being operationally isolated for availability

4

AGILITY

Applications must stay highly decoupled for agility, continuous improvement, and variation

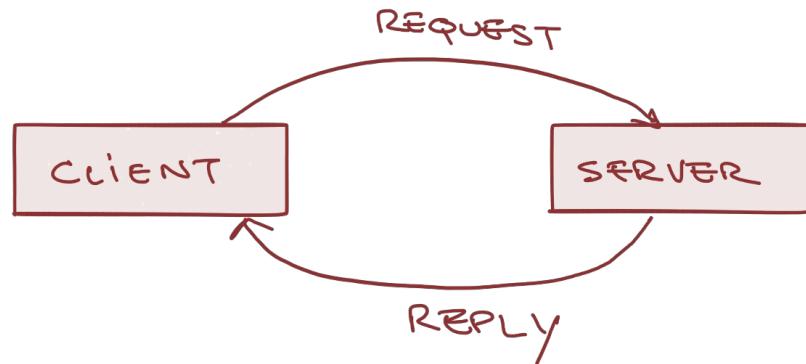
EVENTS

WHAT IS AN EVENT?

- Action or occurrence, something that happened in the past
 - ‘Order created’, ‘user logged in’
- Event characteristics:
 - Immutable
 - Optionally persistent
 - Shareable
- Event types: [1]
 - Notification
 - State Transfer (Command)
 - Event-Sourcing/CQRS

[1] <https://martinfowler.com/articles/201701-event-driven.html>

Request-reply vs. event-driven

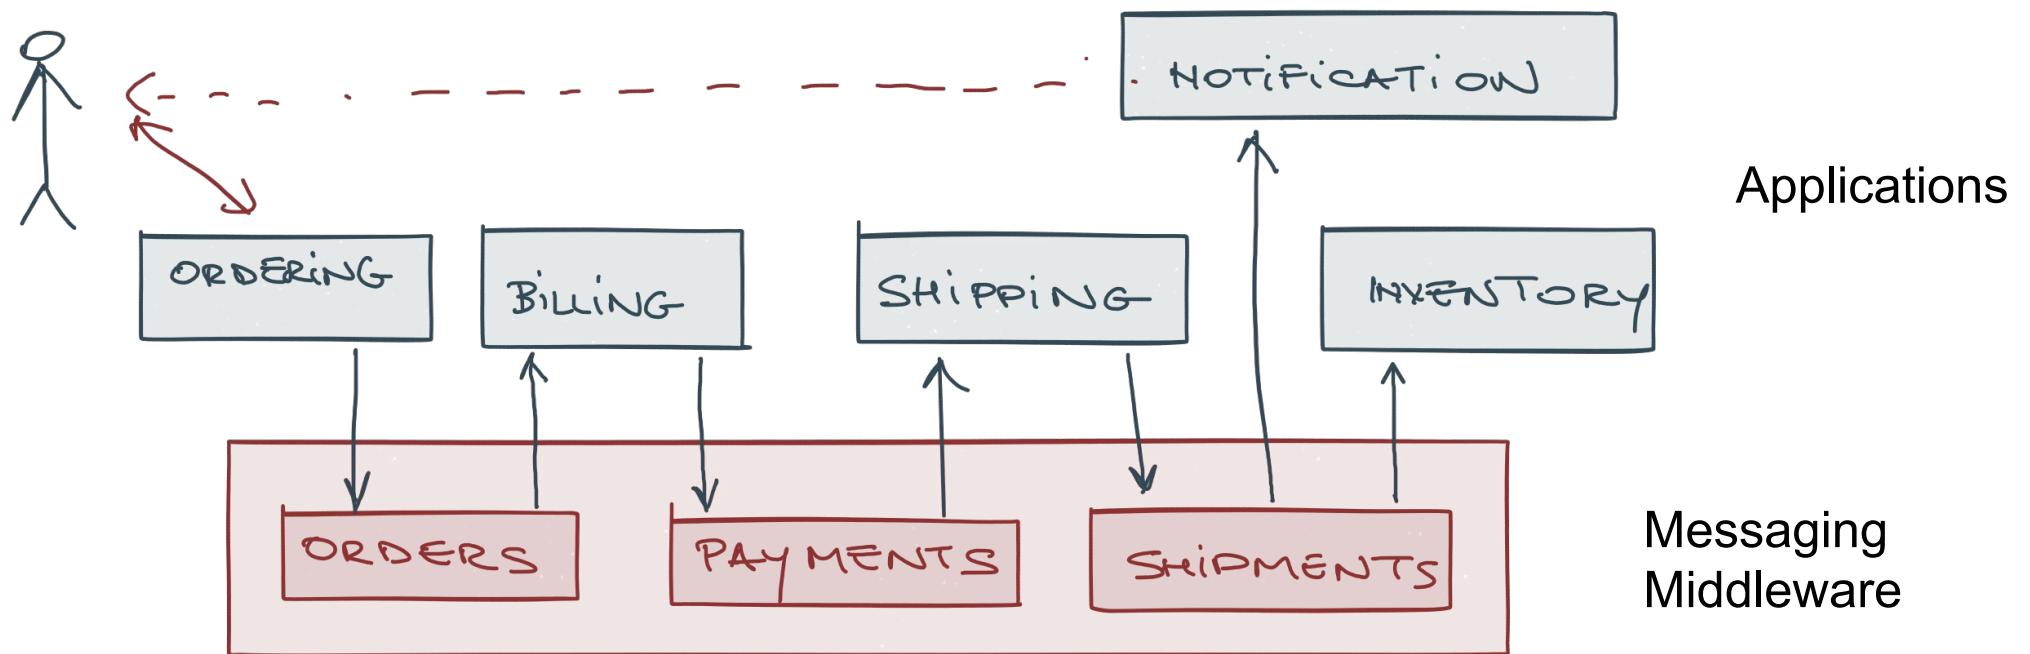


Synchronous & ephemeral
Low composability
Simplified model
Low tolerance to failure
Best practices evolved as REST



Asynchronous and persistent Decoupled
Highly composable
Complex model
High tolerance to failure
Best practices are still evolving

Event-driven microservices



TRADITIONAL MESSAGING



EVENT STREAMING

Advantages

- Store-and-forward
- individual message exchanges (transactionality, acknowledgment, error handling/DLQs), P2P/competing consumer support
- Publish-subscribe support with limitations

Trade-offs

- No replay support
- Requires fast and/or highly available storage infrastructure
- No ordering at scale

Advantages

- long-term persistence, replay, semantic partitioning, large publisher/subscriber imbalances, replay and late-coming subscribers
- Shared nothing data storage model
- Repeatable ordering at scale

Trade-offs

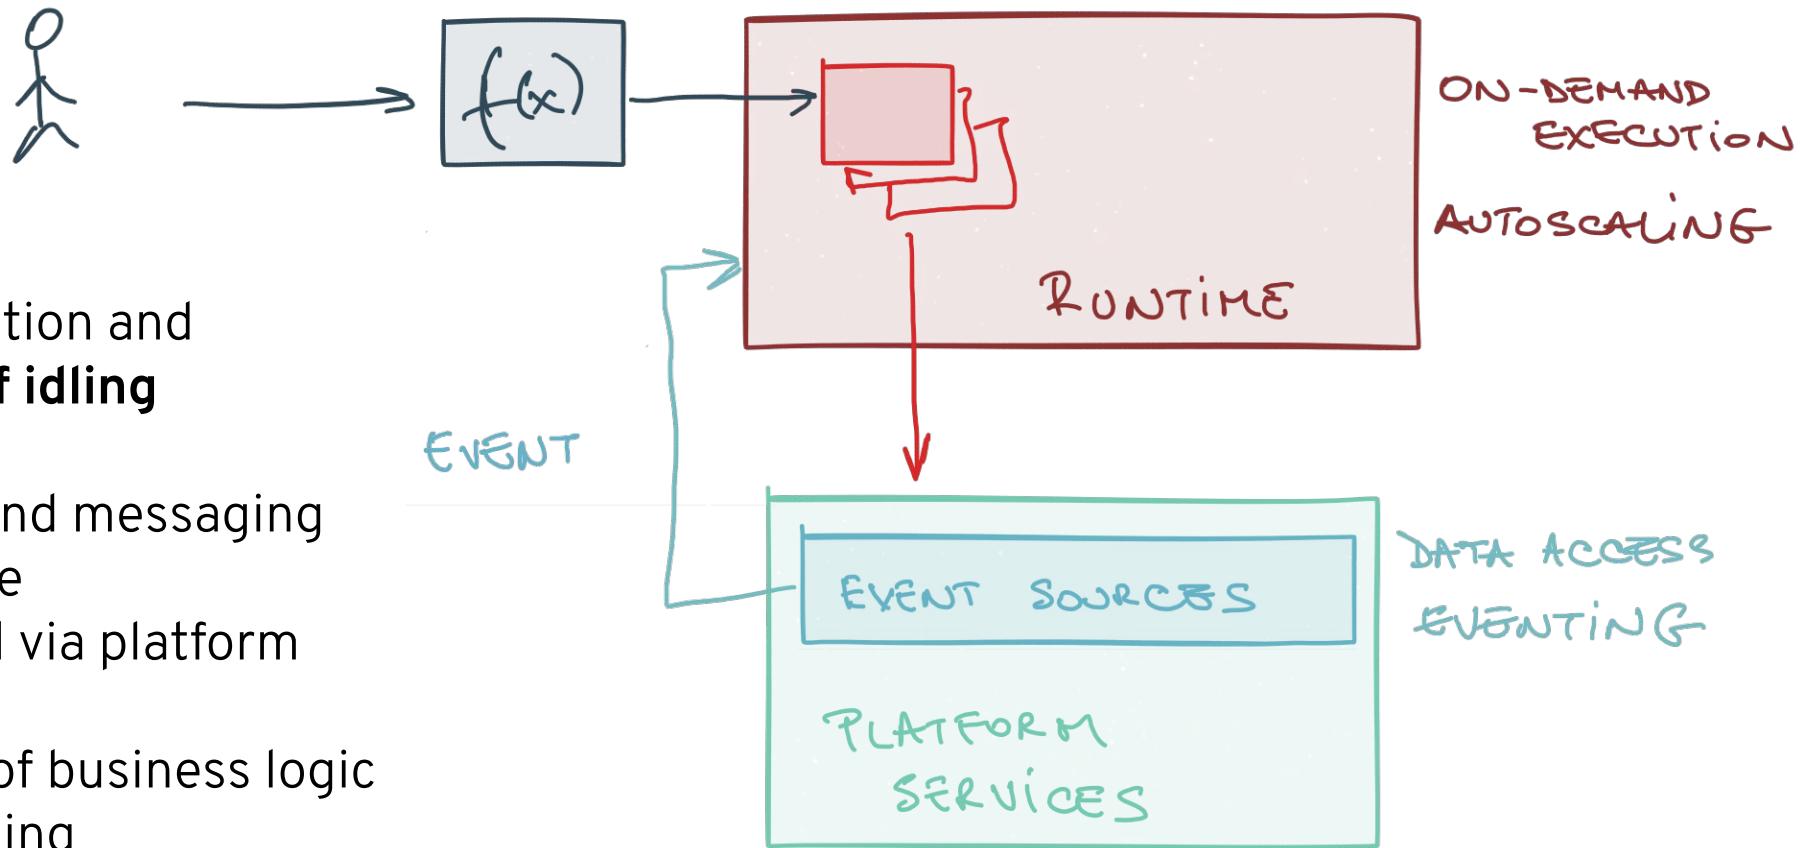
- Weak support for individual message acknowledgment, p2p/competing consumers
- Larger data footprint and extremely fast storage access

Some challenges with microservices ...

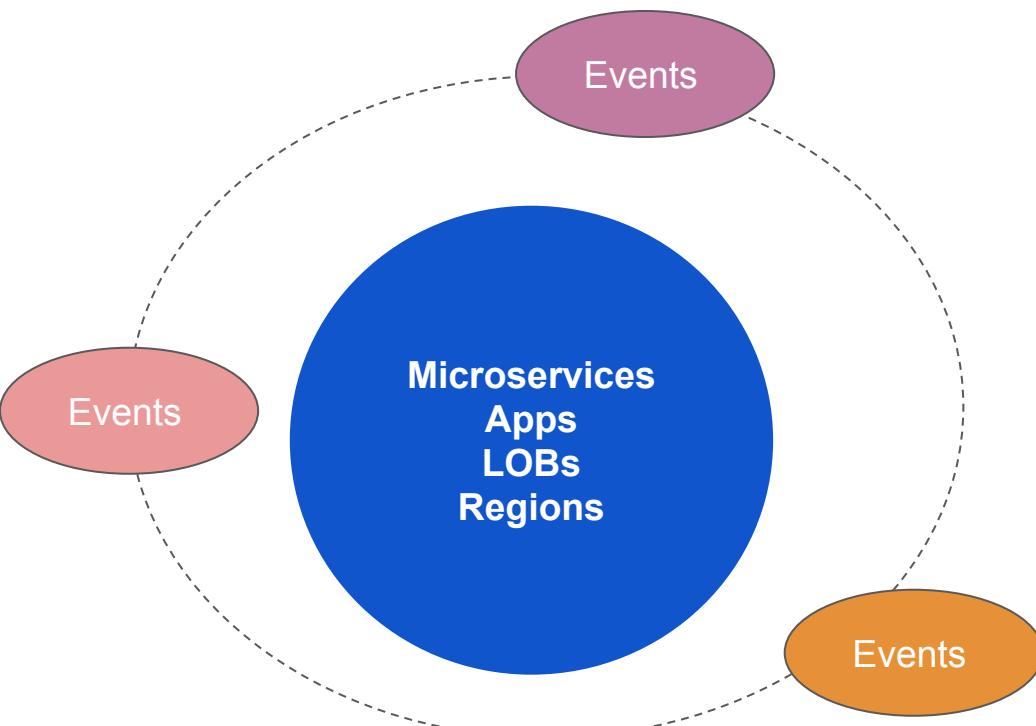
- Utilization
 - Idling under low traffic
 - High resource consumption - memory, disk
- Connectivity
 - Must know broker location
 - Integration with event sources
- Abstraction
 - Must know broker type
 - Dependence on data/payload formats
- Observability
- Security

Serverless model: event-based and elastic

- Elastic execution and **avoidance of idling** (autoscale)
- Utility data and messaging infrastructure
 - Provided via platform services
- **Decoupling** of business logic from messaging infrastructure

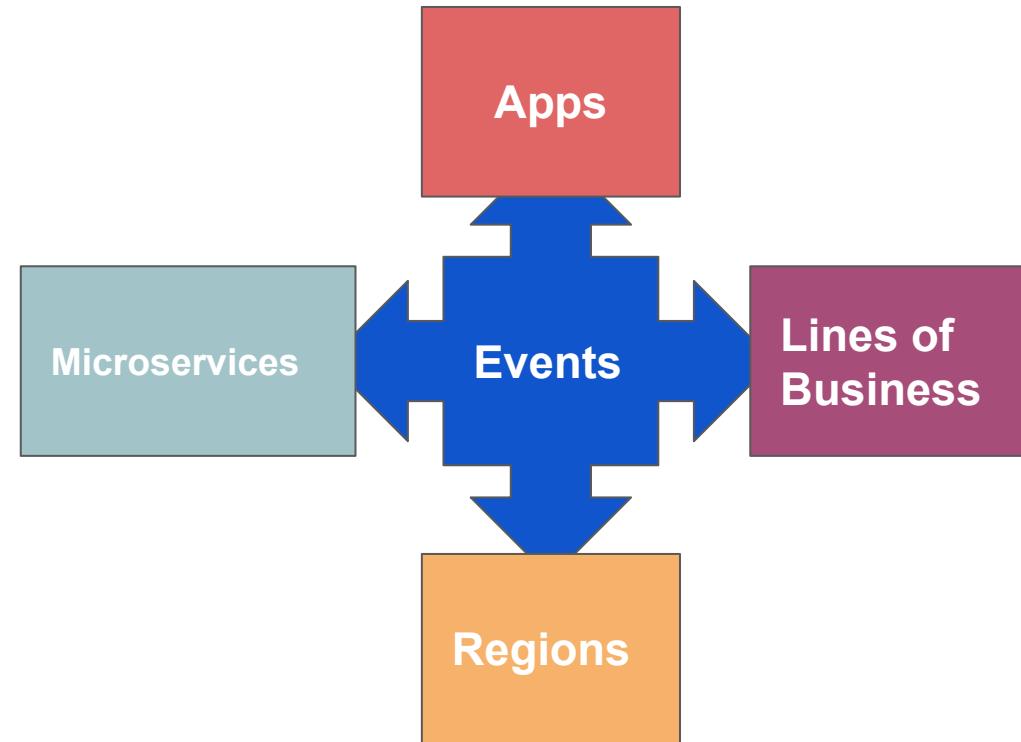


RETHINKING EVENT-DRIVEN ARCHITECTURE



System and data-centric

Events are designed to **respond to ad-hoc connectivity needs**

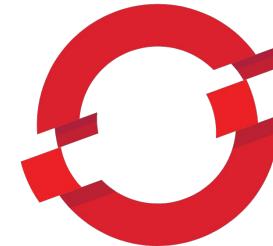


Event-centric

Events are first class citizens that **describe the interactions in the enterprise**

Your technology radar

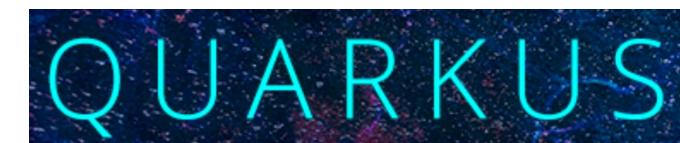
- Service Mesh (e.g. Istio):
 - Provide microservice interconnectivity
- Serverless platforms (e.g. Knative)
 - Container build and on-demand scheduling
- Container-native frameworks (e.g. Quarkus)
 - Optimize Java workloads for containerized apps
- Strimzi - Kafka operator for Kubernetes/OpenShift
- EnMasse - Messaging-as-a-Service for Kubernetes/OpenShift
- FaaS frameworks (e.g. Camel-K)
 - Schedule integration code directly on platform or via Knative



OPENSIFT



Istio



RED HAT AND REACTIVE



MESSAGING BACKBONE



DISTRIBUTE, REACT ON DATA



REACTIVE / FAAS FRAMEWORKS



RULES EVALUATION
COMPLEX EVENTS
AUTOMATION



CAMEL K
REACTIVE INTEGRATION
SERVERLESS



REACTIVE DEVELOPER TOOLING



APP ENVIRONMENT
INFRASTRUCTURE
SERVERLESS / KNATIVE
OPERATOR HUB

GOAL FOR MODULE

In this module you will learn:

- How to adopt event-driven programming and architectures using capabilities of a Red Hat solution
- Adding high performance distributed in-memory cache using Red Hat Data Grid
- Introduce Apache Kafka for responding and processing of application events
- Use Knative Serving and Eventing to begin to move applications to a serverless, event-driven architecture to more efficiently process app events

LAB INSTRUCTIONS

- **Everything is done in browser** - no local commands or installs needed on your laptop
- Tested with **Chrome 75.0.3770.142, Firefox 60.8.0esr**. → **Safari 12.x does not work!**
- If things get weird, just reload browser page
- **Turn off VPN** (we use websockets extensively), **pause AdBlock** for the lab domain (there are no ads)
- To recreate the lab locally, visit
github.com/RedHat-Middleware-Workshops/cloud-native-workshop-v2-infra
- Everyone should have their own **unique logins**, e.g.: user45 / r3dh4t1!

Get Started at [\[LINK\]](#)

Credentials: userXX / r3dh4t1!

If you get stuck, raise hand

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



linkedin.com/company/red-hat



youtube.com/user/RedHatVideos



facebook.com/redhatinc



twitter.com/RedHat